

# **Reflective Journal: AWS MLU Labs**

## **Introduction**

The AWS Machine Learning University (MLU) Labs provided an engaging and hands-on introduction to PyTorch and neural networks. These labs strengthened my understanding of key deep learning concepts, from basic tensor operations to building and training neural networks. This report reflects on my learnings, insights, challenges, and their real-world applications.

## **Lab 01: Getting Started with PyTorch**

### **Summary of Key Learnings**

During this lab, I explored the fundamentals of PyTorch, a widely used deep learning framework. The first notebook introduced PyTorch's tensor operations and automatic differentiation, which are essential for building and optimizing neural networks. I learned how to manipulate data using PyTorch tensors and perform differentiation, a crucial step in backpropagation.

The second notebook focused on constructing a simple neural network using logistic regression. I implemented key components, including data ingestion, model definition, loss computation, and optimization. This helped me understand how neural networks process data and adjust weights to minimize error.

### **Insights & Understanding**

One key insight from this lab was the power of automatic differentiation in PyTorch. Instead of manually computing gradients, PyTorch efficiently tracks operations and calculates derivatives, simplifying the training process. I also gained a better understanding of how tensors function as the backbone of neural networks, enabling flexible and efficient data manipulation.

Another realization was how even a simple logistic regression model can be viewed as a neural network. This reinforced the idea that more complex models, such as deep learning architectures, are built upon foundational principles like matrix multiplications and gradient descent.

### **Challenges Encountered**

One challenge I faced was understanding the role of autograd in PyTorch. Initially, it was unclear how PyTorch computes gradients automatically, but after experimenting with different tensor operations and checking the computational graph, I gained clarity.

Another difficulty was ensuring that the logistic regression model converged correctly. I had to tweak the learning rate and batch size to achieve stable training. This experience highlighted the importance of hyperparameter tuning in deep learning.

### **Application & Relevance**

The concepts learned in this lab have significant applications in AI, particularly in building and optimizing neural networks. Understanding how PyTorch handles tensors and gradients will be invaluable for future projects involving deep learning, such as image classification and natural language processing.

Moreover, the ability to construct and train a neural network from scratch provides a solid foundation for implementing more advanced architectures like convolutional and recurrent neural networks. This knowledge will be crucial as I continue to explore AI applications.

## **Code and Experimentation**

I experimented with different learning rates and batch sizes in the logistic regression model. Lowering the learning rate resulted in slower but more stable convergence, while a higher learning rate led to fluctuations in loss values. This experiment reinforced the importance of choosing appropriate hyperparameters for training stability and efficiency.

Additionally, I modified the dataset size to observe its impact on model performance. With larger datasets, the model achieved better generalization, but training time increased. This highlights the trade-offs involved in training deep learning models.

## **Lab 02: How Neural Networks Learn**

### **Summary of Key Learnings**

This lab provided a hands-on experience in implementing a Multilayer Perceptron (MLP) using PyTorch, emphasizing its application to image classification tasks. Utilizing the Fashion-MNIST dataset, I explored the fundamentals of constructing and training a deep learning model while incorporating dropout layers to improve generalization. The core takeaways included:

- Understanding the architecture of an MLP and its layers (input, hidden, output).
- Leveraging backpropagation and gradient descent for efficient training.
- Implementing dropout layers as a means to prevent overfitting and enhance model robustness.

Through this lab, I developed a deeper appreciation of how deep learning models process and learn from image data and gained confidence in structuring neural networks effectively.

### **Insights & Understanding**

The most profound insight from this lab was recognizing the significance of regularization techniques like dropout in enhancing model performance. Initially, without dropout, the model showed high accuracy on the training set but struggled with generalization. Incorporating dropout layers mitigated this overfitting issue by forcing neurons to work independently rather than relying on co-adaptation.

Another key takeaway was how PyTorch's flexible framework simplifies model development. The ability to construct neural networks using the `torch.nn.Sequential` module streamlined the implementation process. Additionally, I was fascinated by the impact of hyperparameters, particularly how small adjustments in learning rates and batch sizes could significantly influence model training stability and efficiency.

I was also surprised by how well the Fashion-MNIST dataset responded to even a basic MLP model, demonstrating the power of neural networks in extracting meaningful patterns from images. This underscored the importance of dataset complexity in determining model architecture and parameter tuning.

## Challenges Encountered

One of the biggest challenges was tuning hyperparameters such as learning rate, batch size, and dropout probability. Initially, I experimented with a high learning rate, which led to unstable training. Lowering the learning rate improved convergence but prolonged training time. To balance these factors, I adopted the Adam optimizer, which significantly improved both stability and speed.

Another challenge was determining the optimal dropout rate. A high dropout probability (e.g., 0.5) resulted in a significant loss of information, leading to lower training accuracy. Through experimentation, I found that a dropout rate of 0.2-0.3 provided the best balance between preventing overfitting and maintaining model accuracy.

## Application & Relevance

The concepts covered in this lab are highly relevant for various real-world deep learning applications. MLPs form the foundation of advanced architectures such as Convolutional Neural Networks (CNNs) and Transformers, making this knowledge crucial for future AI projects.

Dropout layers, in particular, play an essential role in enhancing model generalization, especially in scenarios where training data is limited or noisy. Moving forward, I plan to incorporate dropout in projects involving text and image classification, ensuring robust performance across diverse datasets.

Additionally, the PyTorch framework demonstrated its capability in building, training, and fine-tuning neural networks efficiently. This experience will be instrumental when experimenting with deeper architectures, transfer learning, and optimizing deep learning pipelines.

## Code and Experimentation

To better understand the effects of hyperparameter changes, I conducted various experiments:

- **Increasing dropout from 0.2 to 0.5** significantly reduced overfitting but also lowered overall accuracy.
- **Decreasing the learning rate from 0.01 to 0.001** improved training stability but required more epochs for convergence.
- **Adding an additional hidden layer** led to a minor boost in accuracy but at the cost of increased computational complexity.

These experiments reinforced the importance of hyperparameter tuning, as minor adjustments can have substantial impacts on model behavior.

## Lab 03: First Example of Neural Networks

### Summary of Key Learnings

This lab provided a comprehensive introduction to building an end-to-end neural network solution for processing text data. Unlike previous labs focusing on image classification, this exercise emphasized handling structured data from the Austin Animal Center Dataset to predict pet adoption outcomes. Key learnings included:

- Importing and preprocessing textual and categorical data for machine learning models.
- Constructing a multilayer neural network in PyTorch.
- Train a model on structured text data and validate its performance.
- Experimenting with different parameters to optimize model accuracy and generalization.

This hands-on experience strengthened my understanding of how deep learning techniques are applied to non-image data and the importance of feature engineering in structured datasets.

### **Insights & Understanding**

A critical takeaway from this lab was the realization that deep learning is not limited to image data but can be highly effective for structured datasets when properly processed. The Austin Animal Center Dataset contained diverse categorical, numerical, and textual features, requiring careful preprocessing before training the model.

I also gained a deeper appreciation for the importance of feature encoding. Handling categorical variables, such as converting text labels into numerical representations through one-hot encoding or embedding layers, significantly impacted model performance.

Additionally, I found the impact of hyperparameter tuning fascinating. Small adjustments in learning rate, batch size, and network architecture dramatically affected model performance, reinforcing the importance of experimentation in neural network training.

### **Challenges Encountered**

One major challenge was efficiently preprocessing the structured dataset. Unlike image data, which is inherently numerical, structured data requires multiple transformations, including handling missing values, encoding categorical variables, and scaling numerical features. To streamline this, I leveraged pandas and PyTorch's data utilities, making data preparation more efficient.

Another challenge was achieving optimal model performance. Initially, the model struggled to generalize, likely due to imbalanced class distributions in the dataset. To address this, I:

- Adjusted the class weights during training to balance predictions.
- Experimented with different activation functions and dropout rates to enhance generalization.
- Utilized cross-validation to validate model performance on unseen data.

### **Application & Relevance**

The ability to apply deep learning to structured data has broad applications beyond image processing. Industries such as healthcare, finance, and retail rely on structured datasets for predictions and decision-making. This lab demonstrated how neural networks can be adapted for such tasks by leveraging appropriate data transformations and optimizations.

Moving forward, I plan to incorporate similar preprocessing techniques and neural network architectures in projects related to text-based data analysis, such as customer behavior prediction, sentiment analysis, and fraud detection.

Moreover, the lab reinforced my understanding of PyTorch's flexibility in handling different data types, making it a versatile tool for machine learning tasks beyond traditional deep learning applications.

## Code and Experimentation

I conducted multiple experiments to optimize model performance:

- **Testing different activation functions:** ReLU performed best in deeper networks, while Tanh provided better stability in shallower models.
- **Adjusting dropout rates:** A dropout of 0.3 improved generalization without significantly impacting training accuracy.
- **Varying batch sizes:** Smaller batches (32-64) led to more stable convergence, whereas larger batches increased training speed but introduced variability in validation accuracy.

These experiments reinforced the importance of iterative optimization and hyperparameter tuning in developing robust machine-learning models.

## Conclusion

Completing these labs has significantly deepened my understanding of PyTorch and neural networks. The hands-on experience solidified theoretical concepts and highlighted practical challenges in training deep learning models. Moving forward, I aim to explore convolutional neural networks (CNNs) and transformers to extend my AI expertise. These labs have provided a strong foundation for building scalable, efficient, and intelligent AI systems.

## References

- PyTorch Documentation: <https://pytorch.org/docs/stable/index.html>
- Deep Learning with PyTorch by Eli Stevens and Thomas Viehmann
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- Austin Animal Center Dataset: <https://www.kaggle.com/datasets/aaronschlegel/austin-animal-center-shelter-intakes-and-outcomes>
- Deep Learning with PyTorch by Eli Stevens, Luca Antiga, and Thomas Viehmann