

✓ Creating Numbers/images with AI: A Hands-on Diffusion Model Exercise

Introduction

In this assignment, you'll learn how to create an AI model that can generate realistic images from scratch using a powerful technique called 'diffusion'. Think of it like teaching AI to draw by first learning how images get blurry and then learning to make them clear again.

What We'll Build

- A diffusion model capable of generating realistic images
- For most students: An AI that generates handwritten digits (0-9) using the MNIST dataset
- For students with more computational resources: Options to work with more complex datasets
- Visual demonstrations of how random noise gradually transforms into clear, recognizable images
- By the end, your AI should create images realistic enough for another AI to recognize them

Dataset Options

This lab offers flexibility based on your available computational resources:

- Standard Option (Free Colab): We'll primarily use the MNIST handwritten digit dataset, which works well with limited GPU memory and completes training in a reasonable time frame. Most examples and code in this notebook are optimized for MNIST.
- Advanced Option: If you have access to more powerful GPUs (either through Colab Pro/Pro+ or your own hardware), you can experiment with more complex datasets like Fashion-MNIST, CIFAR-10, or even face generation. You'll need to adapt the model architecture, hyperparameters, and evaluation metrics accordingly.

Resource Requirements

- Basic MNIST: Works with free Colab GPUs (2-4GB VRAM), ~30 minutes training
- Fashion-MNIST: Similar requirements to MNIST CIFAR-10: Requires more memory (8-12GB VRAM) and longer training (~2 hours)
- Higher resolution images: Requires substantial GPU resources and several hours of training

Before You Start

1. Make sure you're running this in Google Colab or another environment with GPU access
2. Go to 'Runtime' → 'Change runtime type' and select 'GPU' as your hardware accelerator
3. Each code cell has comments explaining what it does
4. Don't worry if you don't understand every detail - focus on the big picture!
5. If working with larger datasets, monitor your GPU memory usage carefully

The concepts you learn with MNIST will scale to more complex datasets, so even if you're using the basic option, you'll gain valuable knowledge about generative AI that applies to more advanced applications.

✓ Step 1: Setting Up Our Tools

First, let's install and import all the tools we need. Run this cell and wait for it to complete.

```

1 # Step 1: Install required packages
2 %pip install einops
3 print("Package installation complete.")
4
5 # Step 2: Import libraries
6 # --- Core PyTorch libraries ---
7 import torch # Main deep learning framework
8 import torch.nn.functional as F # Neural network functions like activation functions
9 import torch.nn as nn # Neural network building blocks (layers)
10 from torch.optim import Adam # Optimization algorithm for training
11
12 # --- Data handling ---
13 from torch.utils.data import Dataset, DataLoader # For organizing and loading our data
14 import torchvision # Library for computer vision datasets and models
15 import torchvision.transforms as transforms # For preprocessing images
16
17 # --- Tensor manipulation ---
18 import random # For random operations
19 from einops.layers.torch import Rearrange # For reshaping tensors in neural networks
20 from einops import rearrange # For elegant tensor reshaping operations
21 import numpy as np # For numerical operations on arrays
22
23 # --- System utilities ---
24 import os # For operating system interactions (used for CPU count)
25
26 # --- Visualization tools ---
27 import matplotlib.pyplot as plt # For plotting images and graphs
28 from PIL import Image # For image processing
29 from torchvision.utils import save_image, make_grid # For saving and displaying image grids
30
31 # Step 3: Set up device (GPU or CPU)
32 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
33 print(f"We'll be using: {device}")
34
35 # Check if we're actually using GPU (for students to verify)
36 if device.type == "cuda":
37     print(f"GPU name: {torch.cuda.get_device_name(0)}")
38     print(f"GPU memory: {torch.cuda.get_device_properties(0).total_memory / 1e9:.2f} GB")
39 else:
40     print("Note: Training will be much slower on CPU. Consider using Google Colab with GPU enabled.")


```

→ Requirement already satisfied: einops in /usr/local/lib/python3.11/dist-packages (0.8.1)
 Package installation complete.
 We'll be using: cuda
 GPU name: Tesla T4
 GPU memory: 15.83 GB

REPRODUCIBILITY AND DEVICE SETUP

```

1 # Step 4: Set random seeds for reproducibility
2 # Diffusion models are sensitive to initialization, so reproducible results help with debugging
3 SEED = 42 # Universal seed value for reproducibility
4 torch.manual_seed(SEED) # PyTorch random number generator
5 np.random.seed(SEED) # NumPy random number generator
6 random.seed(SEED) # Python's built-in random number generator
7


```

```

8 print(f"Random seeds set to {SEED} for reproducible results")
9
10 # Configure CUDA for GPU operations if available
11 if torch.cuda.is_available():
12     torch.cuda.manual_seed(SEED)      # GPU random number generator
13     torch.cuda.manual_seed_all(SEED)  # All GPUs random number generator
14
15     # Ensure deterministic GPU operations
16     # Note: This slightly reduces performance but ensures results are reproducible
17     torch.backends.cudnn.deterministic = True
18     torch.backends.cudnn.benchmark = False
19
20     try:
21         # Check available GPU memory
22         gpu_memory = torch.cuda.get_device_properties(0).total_memory / 1e9 # Convert to GB
23         print(f"Available GPU Memory: {gpu_memory:.1f} GB")
24
25         # Add recommendation based on memory
26         if gpu_memory < 4:
27             print("Warning: Low GPU memory. Consider reducing batch size if you encounter OOM errors.")
28     except Exception as e:
29         print(f"Could not check GPU memory: {e}")
30 else:
31     print("No GPU detected. Training will be much slower on CPU.")
32     print("If you're using Colab, go to Runtime > Change runtime type and select GPU.")

```

Random seeds set to 42 for reproducible results
Available GPU Memory: 15.8 GB

Step 2: Choosing Your Dataset

You have several options for this exercise, depending on your computer's capabilities:

Option 1: MNIST (Basic - Works on Free Colab)

- Content: Handwritten digits (0-9)
- Image size: 28x28 pixels, Grayscale
- Training samples: 60,000
- Memory needed: ~2GB GPU
- Training time: ~15-30 minutes on Colab
- Choose this if:** You're using free Colab or have a basic GPU

Option 2: Fashion-MNIST (Intermediate)

- Content: Clothing items (shirts, shoes, etc.)
- Image size: 28x28 pixels, Grayscale
- Training samples: 60,000
- Memory needed: ~2GB GPU
- Training time: ~15-30 minutes on Colab
- Choose this if:** You want more interesting images but have limited GPU

Option 3: CIFAR-10 (Advanced)

- Content: Real-world objects (cars, animals, etc.)
- Image size: 32x32 pixels, Color (RGB)
- Training samples: 50,000
- Memory needed: ~4GB GPU
- Training time: ~1-2 hours on Colab
- **Choose this if:** You have Colab Pro or a good local GPU (8GB+ memory)

Option 4: CelebA (Expert)

- Content: Celebrity face images
- Image size: 64x64 pixels, Color (RGB)
- Training samples: 200,000
- Memory needed: ~8GB GPU
- Training time: ~3-4 hours on Colab
- **Choose this if:** You have excellent GPU (12GB+ memory)

To use your chosen dataset, uncomment its section in the code below and make sure all others are commented out.

```

1 #=====
2 # SECTION 2: DATASET SELECTION AND CONFIGURATION
3 #=====
4 # STUDENT INSTRUCTIONS:
5 # 1. Choose ONE dataset option based on your available GPU memory
6 # 2. Uncomment ONLY ONE dataset section below
7 # 3. Make sure all other dataset sections remain commented out
8
9 #-----
10 # OPTION 1: MNIST (Basic - 2GB GPU)
11 #-----
12 # Recommended for: Free Colab or basic GPU
13 # Memory needed: ~2GB GPU
14 # Training time: ~15-30 minutes
15 """
16 IMG_SIZE = 28
17 IMG_CH = 1
18 N_CLASSES = 10
19 BATCH_SIZE = 64
20 EPOCHS = 30
21
22 transform = transforms.Compose([
23     transforms.ToTensor(),
24     transforms.Normalize((0.5,), (0.5,)))
25 ])
26
27 # Your code to load the MNIST dataset
28 # Hint: Use torchvision.datasets.MNIST with root='./data', train=True,
29 #        transform=transform, and download=True
30 # Then print a success message
31
32 # Enter your code here:
33 """
34
35 #-----
36 # OPTION 2: Fashion-MNIST (Intermediate - 2GB GPU)

```

```

37 #-----
38 # Uncomment this section to use Fashion-MNIST instead
39
40 IMG_SIZE = 28
41 IMG_CH = 1
42 N_CLASSES = 10
43 BATCH_SIZE = 64
44 EPOCHS = 30
45
46 transform = transforms.Compose([
47     transforms.ToTensor(),
48     transforms.Normalize((0.5,), (0.5,))
49 ])
50
51 # Your code to load the Fashion-MNIST dataset
52 # Hint: Very similar to MNIST but use torchvision.datasets.FashionMNIST
53
54 # Enter your code here:
55 dataset = torchvision.datasets.FashionMNIST(
56     root='./data',
57     train=True,
58     transform=transform,
59     download=True
60 )
61 print("Fashion-MNIST dataset loaded successfully!")
62
63
64 #-----
65 # OPTION 3: CIFAR-10 (Advanced - 4GB+ GPU)
66 #-----
67 # Uncomment this section to use CIFAR-10 instead
68 """
69 IMG_SIZE = 32
70 IMG_CH = 3
71 N_CLASSES = 10
72 BATCH_SIZE = 32 # Reduced batch size for memory
73 EPOCHS = 50 # More epochs for complex data
74
75 # Your code to create the transform and load CIFAR-10
76 # Hint: Use transforms.Normalize with RGB means and stds ((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
77 # Then load torchvision.datasets.CIFAR10
78
79 # Enter your code here:
80
81 """

```

→ 100% |██████████| 26.4M/26.4M [00:02<00:00, 10.5MB/s]
100% |██████████| 29.5k/29.5k [00:00<00:00, 166kB/s]
100% |██████████| 4.42M/4.42M [00:01<00:00, 3.04MB/s]
100% |██████████| 5.15k/5.15k [00:00<00:00, 21.4MB/s]Fashion-MNIST dataset loaded successfully!

```

'\nIMG_SIZE = 32\nIMG_CH = 3\nN_CLASSES = 10\nBATCH_SIZE = 32 # Reduced batch size for memory\nEPOCHS = 50 # More epochs for complex data\n# Your code to create the transform and load CIFAR-10\n# Hint: Use transforms.Normalize with RGB means and stds ((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))\n# Then load torchvision.datasets.CIFAR10\n\n# Enter your code here.\n\n'

```

```

1 #Validating Dataset Selection
2 #Let's add code to validate that a dataset was selected
3 # and check if your GPU has enough memory:
4

```

```

5 # Validate dataset selection
6 if 'dataset' not in locals():
7     raise ValueError("")
8     ✘ ERROR: No dataset selected! Please uncomment exactly one dataset option.
9 Available options:
10 1. MNIST (Basic) - 2GB GPU
11 2. Fashion-MNIST (Intermediate) - 2GB GPU
12 3. CIFAR-10 (Advanced) - 4GB+ GPU
13 4. CelebA (Expert) - 8GB+ GPU
14 """
15
16 # Your code to validate GPU memory requirements
17 # Hint: Check torch.cuda.is_available() and use torch.cuda.get_device_properties(0).total_memory
18 # to get available GPU memory, then compare with dataset requirements
19
20 # Enter your code here:
21 if torch.cuda.is_available():
22     gpu_memory = torch.cuda.get_device_properties(0).total_memory / 1e9
23     required_memory = 2
24     if gpu_memory < required_memory:
25         print(f"Warning: GPU memory ({gpu_memory:.1f} GB) may be insufficient for Fashion-MNIST (requires ~{required_memory} GB)")
26     else:
27         print(f"GPU memory check passed: {gpu_memory:.1f} GB available, {required_memory} GB required")
28 else:
29     print("No GPU available. Fashion-MNIST training will be slower on CPU.")


```

GPU memory check passed: 15.8 GB available, 2 GB required

```

1 #Dataset Properties and Data Loaders
2 #Now let's examine our dataset
3 #and set up the data loaders:
4
5 # Your code to check sample batch properties
6 # Hint: Get a sample batch using next(iter(DataLoader(dataset, batch_size=1)))
7 # Then print information about the dataset shape, type, and value ranges
8
9 # Enter your code here:
10 sample_loader = DataLoader(dataset, batch_size=1, shuffle=True)
11 sample_batch = next(iter(sample_loader))
12 images, labels = sample_batch
13 print(f"Sample batch shape: {images.shape}")
14 print(f"Sample batch type: {images.dtype}")
15 print(f"Value range: min={images.min().item():.2f}, max={images.max().item():.2f}")
16
17 =====
18 # SECTION 3: DATASET SPLITTING AND DATALOADER CONFIGURATION
19 =====
20 # Create train-validation split
21
22 # Your code to create a train-validation split (80% train, 20% validation)
23 # Hint: Use random_split() with appropriate train_size and val_size
24 # Be sure to use a fixed generator for reproducibility
25
26 # Enter your code here:
27 train_size = int(0.8 * len(dataset))
28 val_size = len(dataset) - train_size
29 generator = torch.Generator().manual_seed(SEED)


```

```

1 #!/usr/bin/env python3
2
3 # This script creates dataloaders for training and validation
4 # Hint: Use DataLoader with batch_size=BATCH_SIZE, appropriate shuffle settings,
5 # and num_workers based on available CPU cores
6
7 # Enter your code here:
8 num_workers = min(4, os.cpu_count())
9 train_dataloader = DataLoader(
10     train_dataset,
11     batch_size=BATCH_SIZE,
12     shuffle=True,
13     num_workers=num_workers,
14     pin_memory=True if torch.cuda.is_available() else False
15 )
16
17 val_dataloader = DataLoader(
18     val_dataset,
19     batch_size=BATCH_SIZE,
20     shuffle=False,
21     num_workers=num_workers,
22     pin_memory=True if torch.cuda.is_available() else False
23 )
24
25 print(f"Created train dataloader with {len(train_dataloader)} batches and validation dataloader with {len(val_dataloader)} batches")
26
27 # Sample batch shape: torch.Size([1, 1, 28, 28])
28 # Sample batch type: torch.float32
29 # Value range: min=-1.00, max=1.00
30 # Created train dataloader with 750 batches and validation dataloader with 188 batches

```

Step 3: Building Our Model Components

Now we'll create the building blocks of our AI model. Think of these like LEGO pieces that we'll put together to make our number generator:

- GELUConvBlock: The basic building block that processes images
- DownBlock: Makes images smaller while finding important features
- UpBlock: Makes images bigger again while keeping the important features
- Other blocks: Help the model understand time and what number to generate

```

1 # Basic building block that processes images
2 class GELUConvBlock(nn.Module):
3     def __init__(self, in_ch, out_ch, group_size):
4         """
5             Creates a block with convolution, normalization, and activation
6
7             Args:
8                 in_ch (int): Number of input channels
9                 out_ch (int): Number of output channels
10                group_size (int): Number of groups for GroupNorm
11
12        """
13        super().__init__()
14
15        # Check that group_size is compatible with out_ch
16        if out_ch % group_size != 0:
17            print(f"Warning: out_ch ({out_ch}) is not divisible by group_size ({group_size})")
18
19        # Adjust group_size to be compatible

```

```

18         group_size = min(group_size, out_ch)
19         while out_ch % group_size != 0:
20             group_size -= 1
21         print(f"Adjusted group_size to {group_size}")
22
23     # Your code to create layers for the block
24     # Hint: Use nn.Conv2d, nn.GroupNorm, and nn.GELU activation
25     # Then combine them using nn.Sequential
26
27     # Enter your code here:
28     self.model = nn.Sequential(
29         nn.Conv2d(in_ch, out_ch, kernel_size=3, padding=1),
30         nn.GroupNorm(group_size, out_ch),
31         nn.GELU()
32     )
33 def forward(self, x):
34     # Your code for the forward pass
35     # Hint: Simply pass the input through the model
36
37     # Enter your code here:
38     return self.model(x)

```

```

1 # Rearranges pixels to downsample the image (2x reduction in spatial dimensions)
2 class RearrangePoolBlock(nn.Module):
3     def __init__(self, in_chs, group_size):
4         """
5             Downsamples the spatial dimensions by 2x while preserving information
6
7             Args:
8                 in_chs (int): Number of input channels
9                 group_size (int): Number of groups for GroupNorm
10            """
11     super().__init__()
12
13     # Your code to create the rearrange operation and convolution
14     # Hint: Use Rearrange from einops.layers.torch to reshape pixels
15     # Then add a GELUConvBlock to process the rearranged tensor
16
17     # Enter your code here:
18     self.rearrange = Rearrange('b c (h p1) (w p2) -> b (c p1 p2) h w', p1=2, p2=2)
19     self.conv = GELUConvBlock(in_chs * 4, in_chs, group_size)
20 def forward(self, x):
21     # Your code for the forward pass
22     # Hint: Apply rearrange to downsample, then apply convolution
23
24     # Enter your code here:
25     x = self.rearrange(x)
26     x = self.conv(x)
27     return x

```

```

1 #Let's implement the upsampling block for our U-Net architecture:
2 class DownBlock(nn.Module):
3     """
4         Downsampling block for encoding path in U-Net architecture.
5
6         This block:
7             1. Processes input features with two convolutional blocks

```

```

8     2. Downsamples spatial dimensions by 2x using pixel rearrangement
9
10    Args:
11        in_chs (int): Number of input channels
12        out_chs (int): Number of output channels
13        group_size (int): Number of groups for GroupNorm
14    """
15    def __init__(self, in_chs, out_chs, group_size):
16        super().__init__() # Simplified super() call, equivalent to original
17
18        # Sequential processing of features
19        layers = [
20            GELUConvBlock(in_chs, out_chs, group_size), # First conv block changes channel dimensions
21            GELUConvBlock(out_chs, out_chs, group_size), # Second conv block processes features
22            RearrangePoolBlock(out_chs, group_size) # Downsampling (spatial dims: H,W → H/2,W/2)
23        ]
24        self.model = nn.Sequential(*layers)
25
26        # Log the configuration for debugging
27        print(f"Created DownBlock: in_chs={in_chs}, out_chs={out_chs}, spatial_reduction=2x")
28
29    def forward(self, x):
30        """
31        Forward pass through the DownBlock.
32
33        Args:
34            x (torch.Tensor): Input tensor of shape [B, in_chs, H, W]
35
36        Returns:
37            torch.Tensor: Output tensor of shape [B, out_chs, H/2, W/2]
38        """
39        return self.model(x)

1 #Now let's implement the upsampling block for our U-Net architecture:
2 class UpBlock(nn.Module):
3     """
4     Upsampling block for decoding path in U-Net architecture.
5
6     This block:
7     1. Takes features from the decoding path and corresponding skip connection
8     2. Concatenates them along the channel dimension
9     3. Upsamples spatial dimensions by 2x using transposed convolution
10    4. Processes features through multiple convolutional blocks
11
12    Args:
13        in_chs (int): Number of input channels from the previous layer
14        out_chs (int): Number of output channels
15        group_size (int): Number of groups for GroupNorm
16    """
17    def __init__(self, in_chs, out_chs, group_size):
18        super().__init__()
19
20        # Your code to create the upsampling operation
21        # Hint: Use nn.ConvTranspose2d with kernel_size=2 and stride=2
22        # Note that the input channels will be 2 * in_chs due to concatenation
23
24        # Enter your code here:

```

```

25     self.up = nn.ConvTranspose2d(in_chs, out_chs, kernel_size=2, stride=2)
26     # Your code to create the convolutional blocks
27     # Hint: Use multiple GELUConvBlocks in sequence
28
29     # Enter your code here:
30     self.conv = nn.Sequential(
31         GELUConvBlock(out_chs * 2, out_chs, group_size),
32         GELUConvBlock(out_chs, out_chs, group_size)
33     )
34     # Log the configuration for debugging
35     print(f"Created UpBlock: in_chs={in_chs}, out_chs={out_chs}, spatial_increase=2x")
36
37 def forward(self, x, skip):
38     """
39     Forward pass through the UpBlock.
40
41     Args:
42         x (torch.Tensor): Input tensor from previous layer [B, in_chs, H, W]
43         skip (torch.Tensor): Skip connection tensor from encoder [B, in_chs, 2H, 2W]
44
45     Returns:
46         torch.Tensor: Output tensor with shape [B, out_chs, 2H, 2W]
47     """
48     # Your code for the forward pass
49     # Hint: Concatenate x and skip, then upsample and process
50
51     # Enter your code here:
52     x = self.up(x)
53     x = torch.cat([x, skip], dim=1)
54     x = self.conv(x)
55     return x

```

```

1 # Here we implement the time embedding block for our U-Net architecture:
2 # Helps the model understand time steps in diffusion process
3 class SinusoidalPositionEmbedBlock(nn.Module):
4     """
5     Creates sinusoidal embeddings for time steps in diffusion process.
6
7     This embedding scheme is adapted from the Transformer architecture and
8     provides a unique representation for each time step that preserves
9     relative distance information.
10
11    Args:
12        dim (int): Embedding dimension
13    """
14    def __init__(self, dim):
15        super().__init__()
16        self.dim = dim
17
18    def forward(self, time):
19        """
20        Computes sinusoidal embeddings for given time steps.
21
22        Args:
23            time (torch.Tensor): Time steps tensor of shape [batch_size]
24
25        Returns:

```

```

26         torch.Tensor: Time embeddings of shape [batch_size, dim]
27     """
28     device = time.device
29     half_dim = self.dim // 2
30     embeddings = torch.log(torch.tensor(10000.0, device=device)) / (half_dim - 1)
31     embeddings = torch.exp(torch.arange(half_dim, device=device) * -embeddings)
32     embeddings = time[:, None] * embeddings[None, :]
33     embeddings = torch.cat((embeddings.sin(), embeddings.cos()), dim=-1)
34     return embeddings
35
36

1 # Helps the model understand which number/image to draw (class conditioning)
2 class EmbedBlock(nn.Module):
3     """
4     Creates embeddings for class conditioning in diffusion models.
5
6     This module transforms a one-hot or index representation of a class
7     into a rich embedding that can be added to feature maps.
8
9     Args:
10        input_dim (int): Input dimension (typically number of classes)
11        emb_dim (int): Output embedding dimension
12    """
13    def __init__(self, input_dim, emb_dim):
14        super(EmbedBlock, self).__init__()
15        self.input_dim = input_dim
16
17        # Your code to create the embedding layers
18        # Hint: Use nn.Linear layers with a GELU activation, followed by
19        # nn.Unflatten to reshape for broadcasting with feature maps
20
21        # Enter your code here:
22        self.model = nn.Sequential(
23            nn.Linear(input_dim, emb_dim),
24            nn.GELU(),
25            nn.Linear(emb_dim, emb_dim),
26            nn.Unflatten(-1, (emb_dim, 1, 1))
27        )
28
29    def forward(self, x):
30        """
31        Computes class embeddings for the given class indices.
32
33        Args:
34            x (torch.Tensor): Class indices or one-hot encodings [batch_size, input_dim]
35
36        Returns:
37            torch.Tensor: Class embeddings of shape [batch_size, emb_dim, 1, 1]
38                        (ready to be added to feature maps)
39        """
40        x = x.view(-1, self.input_dim)
41        return self.model(x)
42
43

```

```

1 # Main U-Net model that puts everything together
2 class UNet(nn.Module):
3     """
4         U-Net architecture for diffusion models with time and class conditioning.
5
6         This architecture follows the standard U-Net design with:
7             1. Downsampling path that reduces spatial dimensions
8             2. Middle processing blocks
9             3. Upsampling path that reconstructs spatial dimensions
10            4. Skip connections between symmetric layers
11
12        The model is conditioned on:
13            - Time step (where we are in the diffusion process)
14            - Class labels (what we want to generate)
15
16    Args:
17        T (int): Number of diffusion time steps
18        img_ch (int): Number of image channels
19        img_size (int): Size of input images
20        down_chs (list): Channel dimensions for each level of U-Net
21        t_embed_dim (int): Dimension for time embeddings
22        c_embed_dim (int): Dimension for class embeddings
23    """
24    def __init__(self, T, img_ch, img_size, down_chs, t_embed_dim, c_embed_dim):
25        super().__init__()
26
27        # Your code to create the time embedding
28        # Hint: Use SinusoidalPositionEmbedBlock, nn.Linear, and nn.GELU in sequence
29
30        # Enter your code here:
31        self.time_embed = nn.Sequential(
32            SinusoidalPositionEmbedBlock(t_embed_dim),
33            nn.Linear(t_embed_dim, down_chs[0]),
34            nn.GELU()
35        )
36        # Your code to create the class embedding
37        # Hint: Use the EmbedBlock class you defined earlier
38
39        # Enter your code here:
40        self.class_embed = EmbedBlock(c_embed_dim, down_chs[0])
41        # Your code to create the initial convolution
42        # Hint: Use GELUConvBlock to process the input image
43
44        # Enter your code here:
45        self.init_conv = GELUConvBlock(img_ch, down_chs[0], group_size=8)
46        # Your code to create the downsampling path
47        # Hint: Use nn.ModuleList with DownBlock for each level
48
49        # Enter your code here:
50        self.downs = nn.ModuleList([
51            DownBlock(down_chs[i], down_chs[i+1], group_size=8)
52            for i in range(len(down_chs)-1)
53        ])
54        # Your code to create the middle blocks
55        # Hint: Use GELUConvBlock twice to process features at lowest resolution
56
57        # Enter your code here:
58        self.middle = nn.Sequential(

```

```
59         GELUConvBlock(down_chs[-1], down_chs[-1], group_size=8),
60         GELUConvBlock(down_chs[-1], down_chs[-1], group_size=8)
61     )
62     # Your code to create the upsampling path
63     # Hint: Use nn.ModuleList with UpBlock for each level (in reverse order)
64
65     # Enter your code here:
66     self.ups = nn.ModuleList([
67         UpBlock(down_chs[i], down_chs[i-1], group_size=8)
68         for i in range(len(down_chs)-1, 0, -1)
69     ])
70     # Your code to create the final convolution
71     # Hint: Use nn.Conv2d to project back to the original image channels
72
73     # Enter your code here:
74     self.final_conv = nn.Conv2d(down_chs[0], img_ch, kernel_size=1)
75     print(f"Created UNet with {len(down_chs)} scale levels")
76     print(f"Channel dimensions: {down_chs}")
77
78 def forward(self, x, t, c, c_mask):
79     """
80     Forward pass through the UNet.
81
82     Args:
83         x (torch.Tensor): Input noisy image [B, img_ch, H, W]
84         t (torch.Tensor): Diffusion time steps [B]
85         c (torch.Tensor): Class labels [B, c_embed_dim]
86         c_mask (torch.Tensor): Mask for conditional generation [B, 1]
87
88     Returns:
89         torch.Tensor: Predicted noise in the input image [B, img_ch, H, W]
90     """
91     # Your code for the time embedding
92     # Hint: Process the time steps through the time embedding module
93
94     # Enter your code here:
95     t_emb = self.time_embed(t)
96     t_emb = t_emb.unsqueeze(-1).unsqueeze(-1)
97     # Your code for the class embedding
98     # Hint: Process the class labels through the class embedding module
99
100    # Enter your code here:
101    c_emb = self.class_embed(c)
102    # Your code for the initial feature extraction
103    # Hint: Apply initial convolution to the input
104
105    # Enter your code here:
106    x = self.init_conv(x)
107    x = x + t_emb + c_emb
108    # Your code for the downsampling path and skip connections
109    # Hint: Process the features through each downsampling block
110    # and store the outputs for skip connections
111
112    # Enter your code here:
113    skips = [x]
114    for down in self.downs:
115        x = down(x)
116        skips.append(x)
```

```

117     # Your code for the middle processing and conditioning
118     # Hint: Process features through middle blocks, then add time and class embeddings
119
120     # Enter your code here:
121     x = self.middle(x)
122     # Your code for the upsampling path with skip connections
123     # Hint: Process features through each upsampling block,
124     # combining with corresponding skip connections
125
126     # Enter your code here:
127     skips = skips[::-1][::-1]
128     for up, skip in zip(self.ups, skips):
129         x = up(x, skip)
130     # Your code for the final projection
131     # Hint: Apply the final convolution to get output in image space
132
133     # Enter your code here:
134     x = self.final_conv(x)
135     return x

```

▼ Step 4: Setting Up The Diffusion Process

Now we'll create the process of adding and removing noise from images. Think of it like:

1. Adding fog: Slowly making the image more and more blurry until you can't see it
2. Removing fog: Teaching the AI to gradually make the image clearer
3. Controlling the process: Making sure we can generate specific numbers we want

```

1 # Set up the noise schedule
2 n_steps = 100 # How many steps to go from clear image to noise
3 beta_start = 0.0001 # Starting noise level (small)
4 beta_end = 0.02      # Ending noise level (larger)
5
6 # Create schedule of gradually increasing noise levels
7 beta = torch.linspace(beta_start, beta_end, n_steps).to(device)
8
9 # Calculate important values used in diffusion equations
10 alpha = 1 - beta # Portion of original image to keep at each step
11 alpha_bar = torch.cumprod(alpha, dim=0) # Cumulative product of alphas
12 sqrt_alpha_bar = torch.sqrt(alpha_bar) # For scaling the original image
13 sqrt_one_minus_alpha_bar = torch.sqrt(1 - alpha_bar) # For scaling the noise
14

1
2 # Function to add noise to images (forward diffusion process)
3 def add_noise(x_0, t):
4     """
5     Add noise to images according to the forward diffusion process.
6
7     The formula is: x_t = \sqrt(\alpha_{bar_t}) * x_0 + \sqrt(1-\alpha_{bar_t}) * \epsilon
8     where \epsilon is random noise and \alpha_{bar_t} is the cumulative product of (1-\beta).
9
10    Args:
11        x_0 (torch.Tensor): Original clean image [B, C, H, W]

```

```

12     t (torch.Tensor): Timestep indices indicating noise level [B]
13
14 Returns:
15     tuple: (noisy_image, noise_added)
16         - noisy_image is the image with noise added
17         - noise_added is the actual noise that was added (for training)
18 """
19 # Create random Gaussian noise with same shape as image
20 noise = torch.randn_like(x_0)
21
22 # Get noise schedule values for the specified timesteps
23 # Reshape to allow broadcasting with image dimensions
24 sqrt_alpha_bar_t = sqrt_alpha_bar[t].reshape(-1, 1, 1, 1)
25 sqrt_one_minus_alpha_bar_t = sqrt_one_minus_alpha_bar[t].reshape(-1, 1, 1, 1)
26
27 # Apply the forward diffusion equation:
28 # Mixture of original image (scaled down) and noise (scaled up)      # Your code to apply the forward diffusion equation
29 # Hint: Mix the original image and noise according to the noise schedule
30
31 # Enter your code here:
32 x_t = sqrt_alpha_bar_t * x_0 + sqrt_one_minus_alpha_bar_t * noise
33 return x_t, noise

1 # Function to remove noise from images (reverse diffusion process)
2 @torch.no_grad() # Don't track gradients during sampling (inference only)
3 def remove_noise(x_t, t, model, c, c_mask):
4 """
5 Remove noise from images using the learned reverse diffusion process.
6
7 This implements a single step of the reverse diffusion sampling process.
8 The model predicts the noise in the image, which we then use to partially
9 denoise the image.
10
11 Args:
12     x_t (torch.Tensor): Noisy image at timestep t [B, C, H, W]
13     t (torch.Tensor): Current timestep indices [B]
14     model (nn.Module): U-Net model that predicts noise
15     c (torch.Tensor): Class conditioning (what digit to generate) [B, C]
16     c_mask (torch.Tensor): Mask for conditional generation [B, 1]
17
18 Returns:
19     torch.Tensor: Less noisy image for the next timestep [B, C, H, W]
20 """
21 # Predict the noise in the image using our model
22 predicted_noise = model(x_t, t, c, c_mask)
23
24 # Get noise schedule values for the current timestep
25 alpha_t = alpha[t].reshape(-1, 1, 1, 1)
26 alpha_bar_t = alpha_bar[t].reshape(-1, 1, 1, 1)
27 beta_t = beta[t].reshape(-1, 1, 1, 1)
28
29 # Special case: if we're at the first timestep (t=0), we're done
30 if t[0] == 0:
31     return x_t
32 else:
33     # Calculate the mean of the denoised distribution
34     # This is derived from Bayes' rule and the diffusion process equations

```

```

35     mean = (1 / torch.sqrt(alpha_t)) * (
36         x_t - (beta_t / sqrt_one_minus_alpha_bar[t].reshape(-1, 1, 1, 1)) * predicted_noise
37     )
38
39     # Add a small amount of random noise (variance depends on timestep)
40     # This helps prevent the generation from becoming too deterministic
41     noise = torch.randn_like(x_t)
42
43     # Return the partially denoised image with a bit of new random noise
44     return mean + torch.sqrt(beta_t) * noise

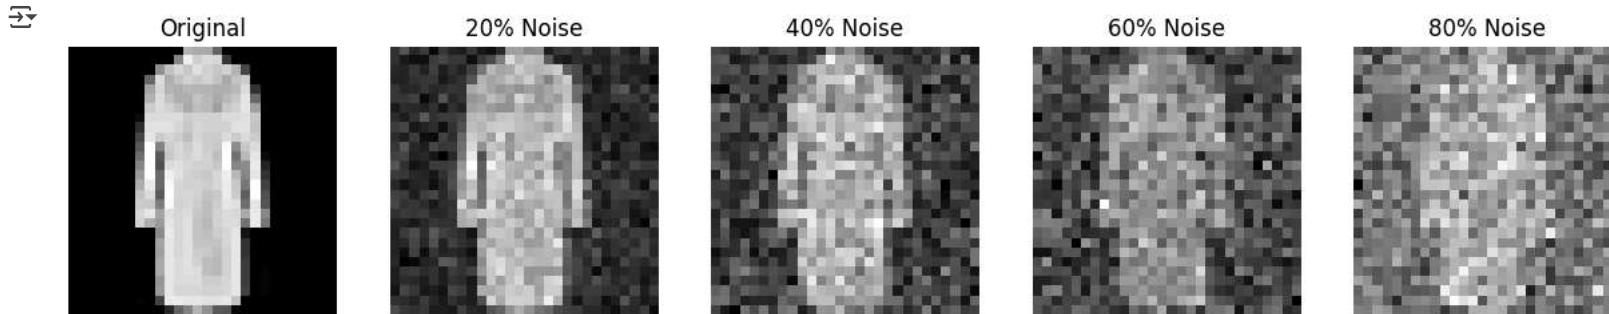
1 # Visualization function to show how noise progressively affects images
2 def show_noise_progression(image, num_steps=5):
3     """
4     Visualize how an image gets progressively noisier in the diffusion process.
5
6     Args:
7         image (torch.Tensor): Original clean image [C, H, W]
8         num_steps (int): Number of noise levels to show
9     """
10    plt.figure(figsize=(15, 3))
11
12    # Show original image
13    plt.subplot(1, num_steps, 1)
14    if IMG_CH == 1: # Grayscale image
15        plt.imshow(image[0].cpu(), cmap='gray')
16    else: # Color image
17        img = image.permute(1, 2, 0).cpu() # Change from [C,H,W] to [H,W,C]
18        if img.min() < 0: # If normalized between -1 and 1
19            img = (img + 1) / 2 # Rescale to [0,1] for display
20        plt.imshow(img)
21    plt.title('Original')
22    plt.axis('off')
23
24    # Show progressively noisier versions
25    for i in range(1, num_steps):
26        # Calculate timestep index based on percentage through the process
27        t_idx = int((i/num_steps) * n_steps)
28        t = torch.tensor([t_idx]).to(device)
29
30        # Add noise corresponding to timestep t
31        noisy_image, _ = add_noise(image.unsqueeze(0), t)
32
33        # Display the noisy image
34        plt.subplot(1, num_steps, i+1)
35        if IMG_CH == 1:
36            plt.imshow(noisy_image[0][0].cpu(), cmap='gray')
37        else:
38            img = noisy_image[0].permute(1, 2, 0).cpu()
39            if img.min() < 0:
40                img = (img + 1) / 2
41            plt.imshow(img)
42        plt.title(f'{int((i/num_steps) * 100)}% Noise')
43        plt.axis('off')
44    plt.show()
45
46 # Show an example of noise progression on a real image

```

```

47 sample_batch = next(iter(train_dataloader)) # Get first batch
48 sample_image = sample_batch[0][0].to(device) # Get first image
49 show_noise_progression(sample_image)
50
51 # Student Activity: Try different noise schedules
52 # Uncomment and modify these lines to experiment:
53 """
54 # Try a non-linear noise schedule
55 beta_alt = torch.linspace(beta_start, beta_end, n_steps)**2
56 alpha_alt = 1 - beta_alt
57 alpha_bar_alt = torch.cumprod(alpha_alt, dim=0)
58 # How would this affect the diffusion process?
59 """

```



```

'\n# Try a non-linear noise schedule\nbeta_alt = torch.linspace(beta_start, beta_end, n_steps)**2\nalpha_alt = 1 - beta_alt\nalpha_bar_alt = torch.cumprod(alpha_alt, dim=0)\n\nHow would this affect the diffusion process?\n'

```

```

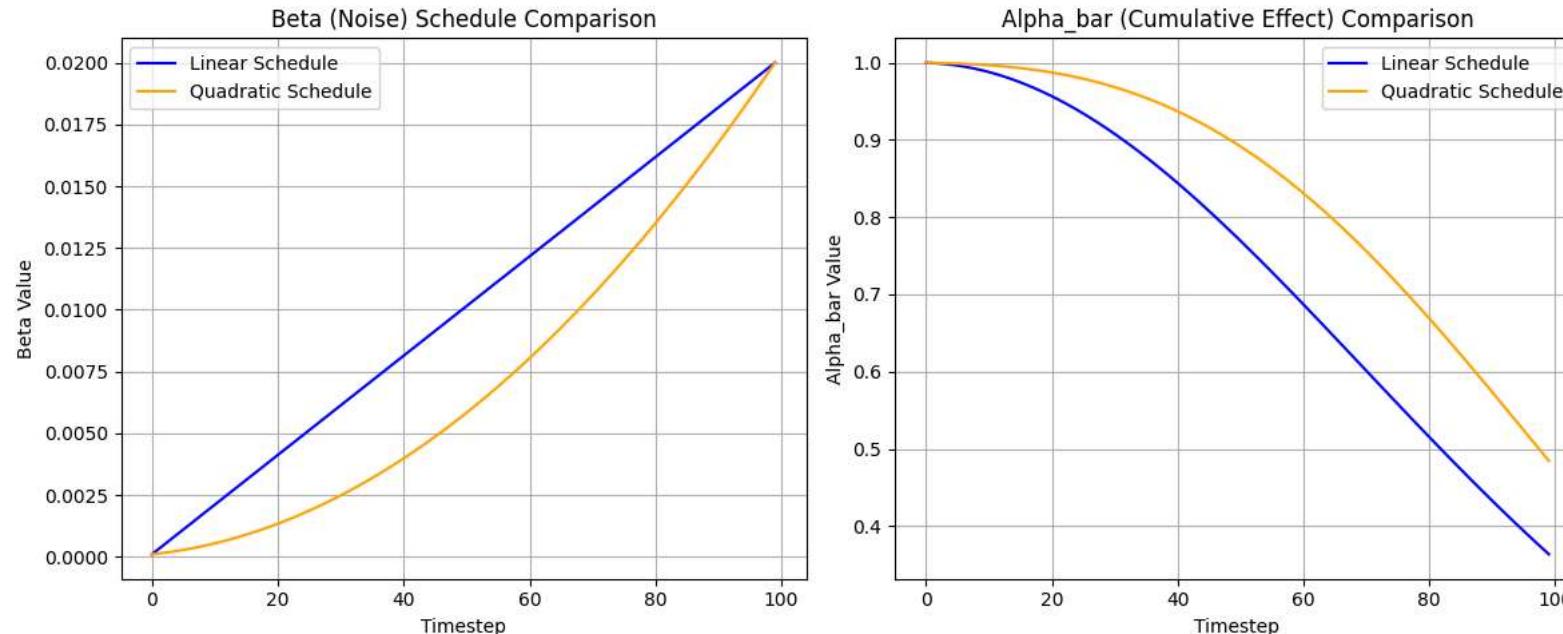
1 # Student Activity: Try different noise schedules
2 print("\nStudent Activity: Experimenting with Different Noise Schedules")
3
4
5 try:
6     n_steps = n_steps
7     beta_start = beta_start
8     beta_end = beta_end
9 except NameError:
10    print("Warning: n_steps, beta_start, and beta_end not defined. Using default values.")
11    n_steps = 1000
12    beta_start = 0.0001
13    beta_end = 0.02
14    print("Warning: Using default values for n_steps, beta_start, and beta_end. Adjust to match your model.")
15
16 # Original linear noise schedule
17 beta_linear = torch.linspace(beta_start, beta_end, n_steps).to(device)
18 alpha_linear = 1 - beta_linear
19 alpha_bar_linear = torch.cumprod(alpha_linear, dim=0)
20
21 # Alternative non-linear noise schedule (quadratic)
22 beta_alt = torch.linspace(beta_start**0.5, beta_end**0.5, n_steps)**2 # Quadratic schedule
23 beta_alt = beta_alt.to(device)
24 alpha_alt = 1 - beta_alt
25 alpha_bar_alt = torch.cumprod(alpha_alt, dim=0)
26
27 # Visualize the noise schedules for comparison
28 plt.figure(figsize=(12, 5))

```

```
29
30 # Plot beta values
31 plt.subplot(1, 2, 1)
32 plt.plot(beta_linear.cpu(), label="Linear Schedule", color='blue')
33 plt.plot(beta_alt.cpu(), label="Quadratic Schedule", color='orange')
34 plt.title("Beta (Noise) Schedule Comparison")
35 plt.xlabel("Timestep")
36 plt.ylabel("Beta Value")
37 plt.legend()
38 plt.grid(True)
39
40 # Plot alpha_bar (cumulative noise effect)
41 plt.subplot(1, 2, 2)
42 plt.plot(alpha_bar_linear.cpu(), label="Linear Schedule", color='blue')
43 plt.plot(alpha_bar_alt.cpu(), label="Quadratic Schedule", color='orange')
44 plt.title("Alpha_bar (Cumulative Effect) Comparison")
45 plt.xlabel("Timestep")
46 plt.ylabel("Alpha_bar Value")
47 plt.legend()
48 plt.grid(True)
49
50 plt.tight_layout()
51 plt.show()
52
53 print("Observations:")
54 print("- Linear Schedule: Noise increases steadily across all timesteps.")
55 print("- Quadratic Schedule: Noise increases more slowly at first, then accelerates.")
56 print("How this affects the diffusion process:")
57 print("- Quadratic schedule may retain more image details early, challenging the model later.")
58
59 # Clean up GPU memory
60 if torch.cuda.is_available():
61     torch.cuda.empty_cache()
```



Student Activity: Experimenting with Different Noise Schedules



Observations:

- Linear Schedule: Noise increases steadily across all timesteps.

- Quadratic Schedule: Noise increases more slowly at first, then accelerates.

How this affects the diffusion process:

- Quadratic schedule may retain more image details early, challenging the model later.

▼ Step 5: Training Our Model

Now we'll teach our AI to generate images. This process:

1. Takes a clear image
2. Adds random noise to it
3. Asks our AI to predict what noise was added
4. Helps our AI learn from its mistakes

This will take a while, but we'll see progress as it learns!

```

1 # Create our model and move it to GPU if available
2 model = UNet(
3     T=n_steps,           # Number of diffusion time steps
4     img_ch=IMG_CH,       # Number of channels in our images (1 for grayscale, 3 for RGB)
5     img_size=IMG_SIZE,   # Size of input images (28 for MNIST, 32 for CIFAR-10)
6     down_chs=(32, 64, 128), # Channel dimensions for each downsampling level
7     t_embed_dim=8,        # Dimension for time step embeddings
8     c_embed_dim=N_CLASSES # Number of classes for conditioning
9 ).to(device)
10

```

```

11 # Print model summary
12 print(f"\n{'='*50}")
13 print("MODEL ARCHITECTURE SUMMARY")
14 print(f"{'='*50}")
15 print(f"Input resolution: {IMG_SIZE}x{IMG_SIZE}")
16 print(f"Input channels: {IMG_CH}")
17 print(f"Time steps: {n_steps}")
18 print(f"Condition classes: {N_CLASSES}")
19 print(f"GPU acceleration: {'Yes' if device.type == 'cuda' else 'No'}")
20
21 # Validate model parameters and estimate memory requirements
22 # Hint: Create functions to count parameters and estimate memory usage
23
24 # Enter your code here:
25 validate_model_parameters(model)
26 # Your code to verify data ranges and integrity
27 # Hint: Create functions to check data ranges in training and validation data
28
29 # Enter your code here:
30 verify_data_range(train_dataloader, "Training")
31 verify_data_range(val_dataloader, "Validation")
32 # Set up the optimizer with parameters tuned for diffusion models
33 # Note: Lower learning rates tend to work better for diffusion models
34 initial_lr = 0.001 # Starting learning rate
35 weight_decay = 1e-5 # L2 regularization to prevent overfitting
36
37 optimizer = Adam(
38     model.parameters(),
39     lr=initial_lr,
40     weight_decay=weight_decay
41 )
42
43 # Learning rate scheduler to reduce LR when validation loss plateaus
44 # This helps fine-tune the model toward the end of training
45 scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
46     optimizer,
47     mode='min',           # Reduce LR when monitored value stops decreasing
48     factor=0.5,          # Multiply LR by this factor
49     patience=5,          # Number of epochs with no improvement after which LR will be reduced
50     verbose=True,         # Print message when LR is reduced
51     min_lr=1e-6           # Lower bound on the learning rate
52 )
53
54 # STUDENT EXPERIMENT:
55 # Try different channel configurations and see how they affect:
56 # 1. Model size (parameter count)
57 # 2. Training time
58 # 3. Generated image quality
59 #
60 # Suggestions:
61 # - Smaller: down_chs=(16, 32, 64)
62 # - Larger: down_chs=(64, 128, 256, 512)

→ Created DownBlock: in_chs=32, out_chs=64, spatial_reduction=2x
Created DownBlock: in_chs=64, out_chs=128, spatial_reduction=2x
Created UpBlock: in_chs=128, out_chs=64, spatial_increase=2x
Created UpBlock: in_chs=64, out_chs=32, spatial_increase=2x
Created UNet with 3 scale levels
Channel dimensions: (32, 64, 128)

```

```
=====
MODEL ARCHITECTURE SUMMARY
=====
Input resolution: 28x28
Input channels: 1
Time steps: 100
Condition classes: 10
GPU acceleration: Yes
Total parameters: 1,493,153
Trainable parameters: 1,493,153
Estimated GPU memory usage: 22.8 MB

Training range check:
Shape: torch.Size([64, 1, 28, 28])
Data type: torch.float32
Min value: -1.00
Max value: 1.00
Contains NaN: False
Contains Inf: False

Validation range check:
Shape: torch.Size([64, 1, 28, 28])
Data type: torch.float32
Min value: -1.00
Max value: 1.00
Contains NaN: False
Contains Inf: False
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
  warnings.warn(
  )

1 # Define helper functions needed for training and evaluation
2 def validate_model_parameters(model):
3     """
4     Counts model parameters and estimates memory usage.
5     """
6     total_params = sum(p.numel() for p in model.parameters())
7     trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
8
9     print(f"Total parameters: {total_params:,}")
10    print(f"Trainable parameters: {trainable_params:,}")
11
12    # Estimate memory requirements (very approximate)
13    param_memory = total_params * 4 / (1024 ** 2) # MB for params (float32)
14    grad_memory = trainable_params * 4 / (1024 ** 2) # MB for gradients
15    buffer_memory = param_memory * 2 # Optimizer state, forward activations, etc.
16
17    print(f"Estimated GPU memory usage: {param_memory + grad_memory + buffer_memory:.1f} MB")
18
19 # Define helper functions for verifying data ranges
20 def verify_data_range(dataloader, name="Dataset"):
21     """
22     Verifies the range and integrity of the data.
23     """
24     batch = next(iter(dataloader))[0]
25     print(f"\n{name} range check:")
26     print(f"Shape: {batch.shape}")
27     print(f"Data type: {batch.dtype}")
28     print(f"Min value: {batch.min().item():.2f}")
29     print(f"Max value: {batch.max().item():.2f}")
```

```
30     print(f"Contains NaN: {torch.isnan(batch).any().item()}")
31     print(f"Contains Inf: {torch.isinf(batch).any().item()}")
32
33 # Define helper functions for generating samples during training
34 def generate_samples(model, n_samples=10):
35     """
36     Generates sample images using the model for visualization during training.
37     """
38     model.eval()
39     with torch.no_grad():
40         # Generate digits 0-9 for visualization
41         samples = []
42         for digit in range(min(n_samples, 10)):
43             # Start with random noise
44             x = torch.randn(1, IMG_CH, IMG_SIZE, IMG_SIZE).to(device)
45
46             # Set up conditioning for the digit
47             c = torch.tensor([digit]).to(device)
48             c_one_hot = F.one_hot(c, N_CLASSES).float().to(device)
49             c_mask = torch.ones_like(c.unsqueeze(-1)).to(device)
50
51             # Remove noise step by step
52             for t in range(n_steps-1, -1, -1):
53                 t_batch = torch.full((1,), t).to(device)
54                 x = remove_noise(x, t_batch, model, c_one_hot, c_mask)
55
56             samples.append(x)
57
58         # Combine samples and display
59         samples = torch.cat(samples, dim=0)
60         grid = make_grid(samples, nrow=min(n_samples, 5), normalize=True)
61
62         plt.figure(figsize=(10, 4))
63
64         # Display based on channel configuration
65         if IMG_CH == 1:
66             plt.imshow(grid[0].cpu(), cmap='gray')
67         else:
68             plt.imshow(grid.permute(1, 2, 0).cpu())
69
70         plt.axis('off')
71         plt.title('Generated Samples')
72         plt.show()
73
74 # Define helper functions for safely saving models
75 def safe_save_model(model, path, optimizer=None, epoch=None, best_loss=None):
76     """
77     Safely saves model with error handling and backup.
78     """
79     try:
80         # Create a dictionary with all the elements to save
81         save_dict = {
82             'model_state_dict': model.state_dict(),
83         }
84
85         # Add optional elements if provided
86         if optimizer is not None:
87             save_dict['optimizer_state_dict'] = optimizer.state_dict()

```

```

88     if epoch is not None:
89         save_dict['epoch'] = epoch
90     if best_loss is not None:
91         save_dict['best_loss'] = best_loss
92
93     # Create a backup of previous checkpoint if it exists
94     if os.path.exists(path):
95         backup_path = path + '.backup'
96         try:
97             os.replace(path, backup_path)
98             print(f"Created backup at {backup_path}")
99         except Exception as e:
100             print(f"Warning: Could not create backup - {e}")
101
102     # Save the new checkpoint
103     torch.save(save_dict, path)
104     print(f"Model successfully saved to {path}")
105
106 except Exception as e:
107     print(f"Error saving model: {e}")
108     print("Attempting emergency save...")
109
110     try:
111         emergency_path = path + '.emergency'
112         torch.save(model.state_dict(), emergency_path)
113         print(f"Emergency save successful: {emergency_path}")
114     except:
115         print("Emergency save failed. Could not save model.")

1 # Implementation of the training step function
2 def train_step(x, c):
3     """
4     Performs a single training step for the diffusion model.
5
6     This function:
7     1. Prepares class conditioning
8     2. Samples random timesteps for each image
9     3. Adds corresponding noise to the images
10    4. Asks the model to predict the noise
11    5. Calculates the loss between predicted and actual noise
12
13    Args:
14        x (torch.Tensor): Batch of clean images [batch_size, channels, height, width]
15        c (torch.Tensor): Batch of class labels [batch_size]
16
17    Returns:
18        torch.Tensor: Mean squared error loss value
19    """
20    # Convert number labels to one-hot encoding for class conditioning
21    # Example: Label 3 -> [0, 0, 0, 1, 0, 0, 0, 0, 0] for MNIST
22    c_one_hot = F.one_hot(c, N_CLASSES).float().to(device)
23
24    # Create conditioning mask (all ones for standard training)
25    # This would be used for classifier-free guidance if implemented
26    c_mask = torch.ones_like(c.unsqueeze(-1)).to(device)
27
28    # Pick random timesteps for each image in the batch

```

```

29 # Different timesteps allow the model to learn the entire diffusion process
30 t = torch.randint(0, n_steps, (x.shape[0],)).to(device)
31
32 # Add noise to images according to the forward diffusion process
33 # This simulates images at different stages of the diffusion process
34 # Hint: Use the add_noise function you defined earlier
35
36 # Enter your code here:
37 x_t, noise = add_noise(x, t)
38 # The model tries to predict the exact noise that was added
39 # This is the core learning objective of diffusion models
40 predicted_noise = model(x_t, t, c_one_hot, c_mask)
41
42 # Calculate loss: how accurately did the model predict the noise?
43 # MSE loss works well for image-based diffusion models
44 # Hint: Use F.mse_loss to compare predicted and actual noise
45
46 # Enter your code here:
47 loss = F.mse_loss(predicted_noise, noise)
48 return loss

1 # Implementation of the main training loop
2 # Training configuration
3 early_stopping_patience = 10 # Number of epochs without improvement before stopping
4 gradient_clip_value = 1.0 # Maximum gradient norm for stability
5 display_frequency = 100 # How often to show progress (in steps)
6 generate_frequency = 500 # How often to generate samples (in steps)
7
8 # Progress tracking variables
9 best_loss = float('inf')
10 train_losses = []
11 val_losses = []
12 no_improve_epochs = 0
13
14 # Training loop
15 print("\n" + "="*50)
16 print("STARTING TRAINING")
17 print("="*50)
18 model.train()
19
20 # Wrap the training loop in a try-except block for better error handling:
21 # Your code for the training loop
22 # Hint: Use a try-except block for better error handling
23 # Process each epoch and each batch, with validation after each epoch
24
25 # Enter your code here:
26 try:
27     for epoch in range(EPOCHS):
28         print(f"\nEpoch {epoch+1}/{EPOCHS}")
29         print("-" * 20)
30
31         # Training phase
32         model.train()
33         epoch_losses = []
34
35         # Process each batch
36         for step, (images, labels) in enumerate(train_dataloader): # Fixed: dataloader → train_dataloader

```

```
37     images = images.to(device)
38     labels = labels.to(device)
39
40     # Training step
41     optimizer.zero_grad()
42     loss = train_step(images, labels)
43     loss.backward()
44
45     # Add gradient clipping for stability
46     torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=gradient_clip_value)
47
48     optimizer.step()
49     epoch_losses.append(loss.item())
50
51     # Show progress at regular intervals
52     if step % display_frequency == 0:
53         print(f" Step {step}/{len(train_dataloader)}, Loss: {loss.item():.4f}")
54
55     # Generate samples less frequently to save time
56     if step % generate_frequency == 0 and step > 0:
57         print(" Generating samples...")
58         generate_samples(model, n_samples=5)
59
60     # End of epoch - calculate average training loss
61     avg_train_loss = sum(epoch_losses) / len(epoch_losses)
62     train_losses.append(avg_train_loss)
63     print(f"\nTraining - Epoch {epoch+1} average loss: {avg_train_loss:.4f}")
64
65     # Validation phase
66     model.eval()
67     val_epoch_losses = []
68     print("Running validation...")
69
70     with torch.no_grad(): # Disable gradients for validation
71         for val_images, val_labels in val_dataloader:
72             val_images = val_images.to(device)
73             val_labels = val_labels.to(device)
74
75             # Calculate validation loss
76             val_loss = train_step(val_images, val_labels)
77             val_epoch_losses.append(val_loss.item())
78
79     # Calculate average validation loss
80     avg_val_loss = sum(val_epoch_losses) / len(val_epoch_losses)
81     val_losses.append(avg_val_loss)
82     print(f"Validation - Epoch {epoch+1} average loss: {avg_val_loss:.4f}")
83
84     # Learning rate scheduling based on validation loss
85     scheduler.step(avg_val_loss)
86     current_lr = optimizer.param_groups[0]['lr']
87     print(f"Learning rate: {current_lr:.6f}")
88
89     # Generate samples at the end of each epoch
90     if epoch % 2 == 0 or epoch == EPOCHS - 1:
91         print("\nGenerating samples for visual progress check...")
92         generate_samples(model, n_samples=10)
93
94     # Save best model based on validation loss
```

```
95     if avg_val_loss < best_loss:
96         best_loss = avg_val_loss
97         # Use safe_save_model instead of just saving state_dict
98         safe_save_model(model, 'best_diffusion_model.pt', optimizer, epoch, best_loss)
99         print(f"✓ New best model saved! (Val Loss: {best_loss:.4f})")
100        no_improve_epochs = 0
101    else:
102        no_improve_epochs += 1
103        print(f"No improvement for {no_improve_epochs}/{early_stopping_patience} epochs")
104
105    # Early stopping
106    if no_improve_epochs >= early_stopping_patience:
107        print("\nEarly stopping triggered! No improvement in validation loss.")
108        break
109
110    # Plot loss curves every few epochs
111    if epoch % 5 == 0 or epoch == EPOCHS - 1:
112        plt.figure(figsize=(10, 5))
113        plt.plot(train_losses, label='Training Loss')
114        plt.plot(val_losses, label='Validation Loss')
115        plt.xlabel('Epoch')
116        plt.ylabel('Loss')
117        plt.title('Training and Validation Loss')
118        plt.legend()
119        plt.grid(True)
120        plt.show()
121
122 except Exception as e:
123     print(f"\nTraining interrupted due to error: {e}")
124     safe_save_model(model, 'emergency_save.pt', optimizer, epoch, best_loss)
125     raise
126
127 # Final wrap-up
128 print("\n" + "="*50)
129 print("TRAINING COMPLETE")
130 print("="*50)
131 print(f"Best validation loss: {best_loss:.4f}")
132
133 # Generate final samples
134 print("Generating final samples...")
135 generate_samples(model, n_samples=10)
136
137 # Display final loss curves
138 plt.figure(figsize=(12, 5))
139 plt.plot(train_losses, label='Training Loss')
140 plt.plot(val_losses, label='Validation Loss')
141 plt.xlabel('Epoch')
142 plt.ylabel('Loss')
143 plt.title('Training and Validation Loss')
144 plt.legend()
145 plt.grid(True)
146 plt.show()
147
148 # Clean up memory
149 torch.cuda.empty_cache()
```

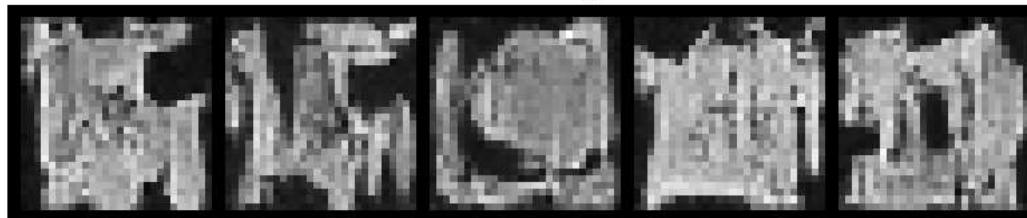


```
=====
STARTING TRAINING
=====
```

Epoch 1/30

```
-----
Step 0/750, Loss: 1.0427
Step 100/750, Loss: 0.1921
Step 200/750, Loss: 0.1574
Step 300/750, Loss: 0.1777
Step 400/750, Loss: 0.1601
Step 500/750, Loss: 0.1775
Generating samples...
```

Generated Samples



```
Step 600/750, Loss: 0.1207
Step 700/750, Loss: 0.1343
```

Training - Epoch 1 average loss: 0.1706

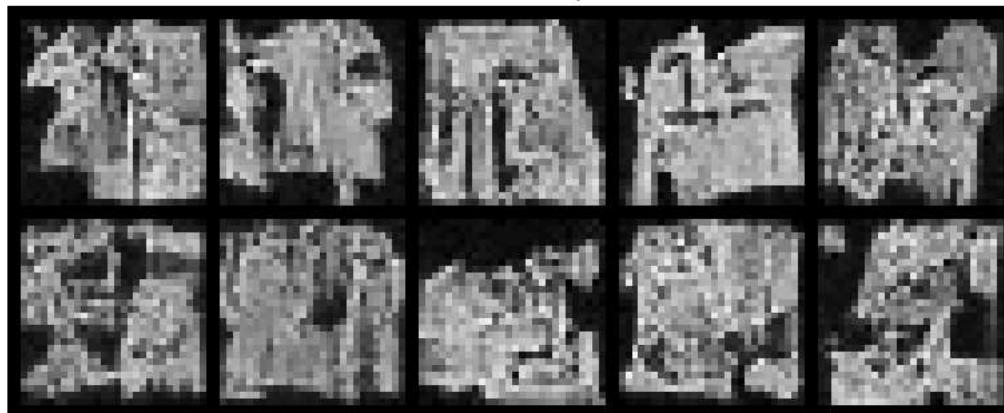
Running validation...

Validation - Epoch 1 average loss: 0.1370

Learning rate: 0.001000

Generating samples for visual progress check...

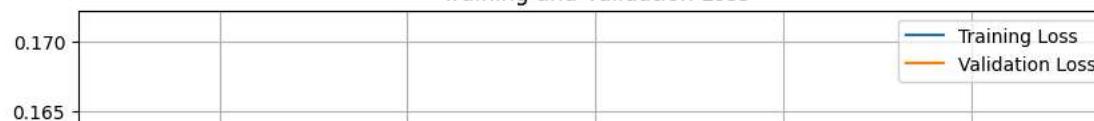
Generated Samples

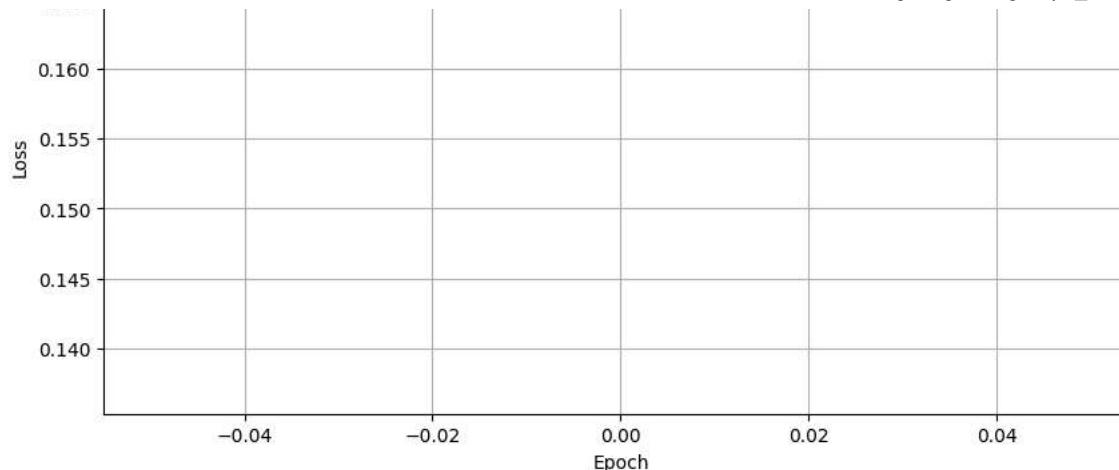


Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.1370)

Training and Validation Loss

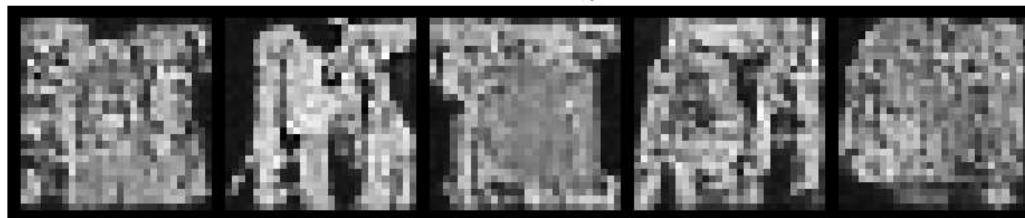




Epoch 2/30

```
-----  
Step 0/750, Loss: 0.1155  
Step 100/750, Loss: 0.1262  
Step 200/750, Loss: 0.1168  
Step 300/750, Loss: 0.0997  
Step 400/750, Loss: 0.1211  
Step 500/750, Loss: 0.1323  
Generating samples...
```

Generated Samples



```
Step 600/750, Loss: 0.1330  
Step 700/750, Loss: 0.1309
```

Training - Epoch 2 average loss: 0.1289

Running validation...

Validation - Epoch 2 average loss: 0.1258

Learning rate: 0.001000

Created backup at best_diffusion_model.pt.backup

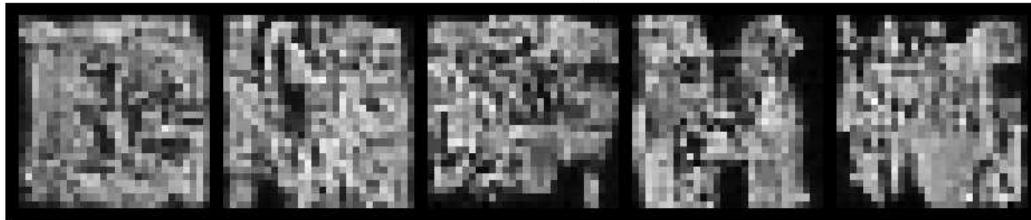
Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.1258)

Epoch 3/30

```
-----  
Step 0/750, Loss: 0.1142  
Step 100/750, Loss: 0.1627  
Step 200/750, Loss: 0.1011  
Step 300/750, Loss: 0.1317  
Step 400/750, Loss: 0.1330  
Step 500/750, Loss: 0.1116  
Generating samples...
```

Generated Samples



Step 600/750, Loss: 0.1440
Step 700/750, Loss: 0.1558

Training - Epoch 3 average loss: 0.1215

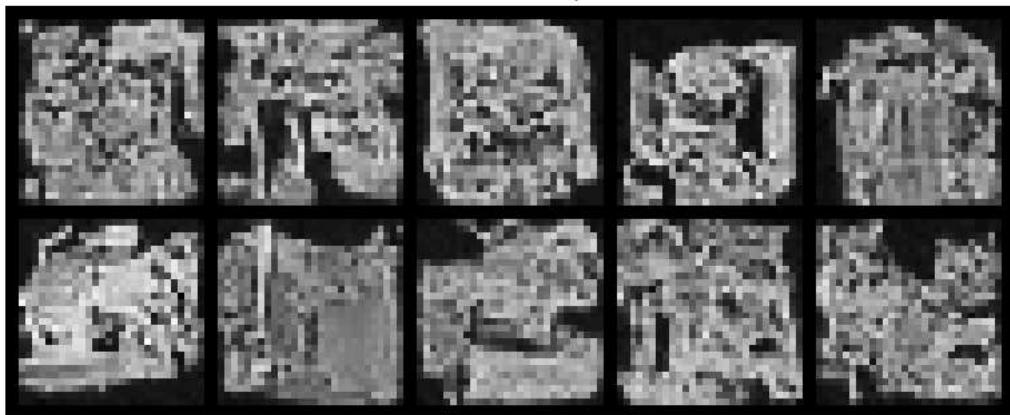
Running validation...

Validation - Epoch 3 average loss: 0.1209

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



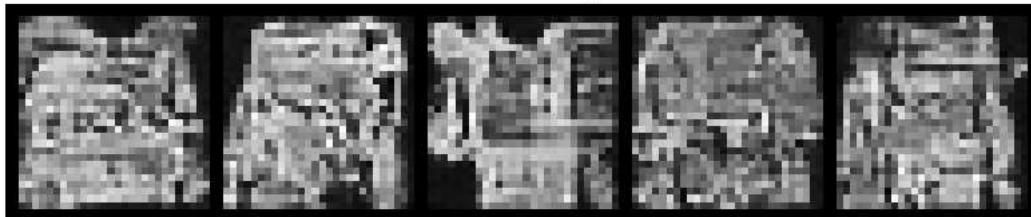
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1209)

Epoch 4/30

Step 0/750, Loss: 0.1228
Step 100/750, Loss: 0.1257
Step 200/750, Loss: 0.1245
Step 300/750, Loss: 0.1306
Step 400/750, Loss: 0.1097
Step 500/750, Loss: 0.1181

Generating samples...

Generated Samples



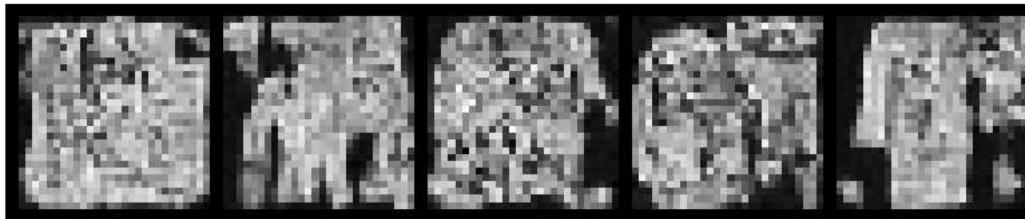
```
Step 600/750, Loss: 0.1057
Step 700/750, Loss: 0.0981
```

```
Training - Epoch 4 average loss: 0.1184
Running validation...
Validation - Epoch 4 average loss: 0.1182
Learning rate: 0.001000
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1182)
```

Epoch 5/30

```
-----  
Step 0/750, Loss: 0.1020
Step 100/750, Loss: 0.0994
Step 200/750, Loss: 0.1365
Step 300/750, Loss: 0.1183
Step 400/750, Loss: 0.1022
Step 500/750, Loss: 0.1055
Generating samples...
```

Generated Samples

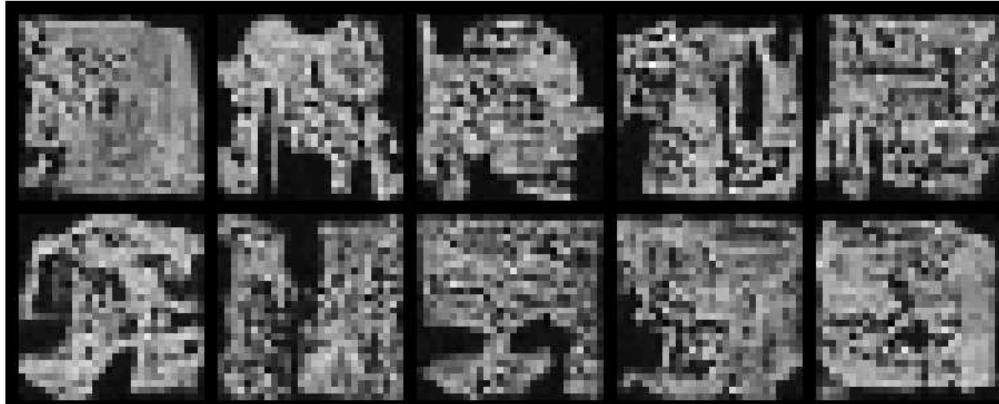


```
Step 600/750, Loss: 0.0946
Step 700/750, Loss: 0.1142
```

```
Training - Epoch 5 average loss: 0.1163
Running validation...
Validation - Epoch 5 average loss: 0.1147
Learning rate: 0.001000
```

Generating samples for visual progress check...

Generated Samples

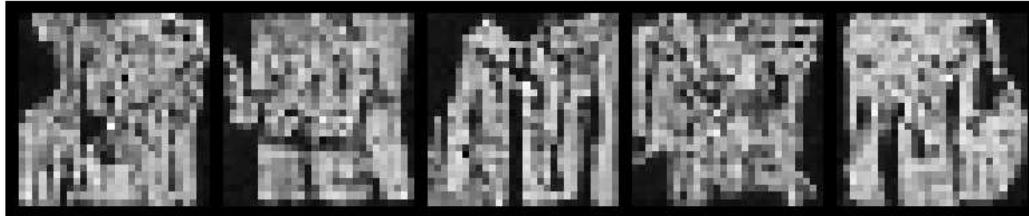


```
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1147)
```

Epoch 6/30

```
-----  
Step 0/750, Loss: 0.1040  
Step 100/750, Loss: 0.1268  
Step 200/750, Loss: 0.1264  
Step 300/750, Loss: 0.1126  
Step 400/750, Loss: 0.1247  
Step 500/750, Loss: 0.1412  
Generating samples...
```

Generated Samples



```
Step 600/750, Loss: 0.1103  
Step 700/750, Loss: 0.1240
```

Training - Epoch 6 average loss: 0.1151

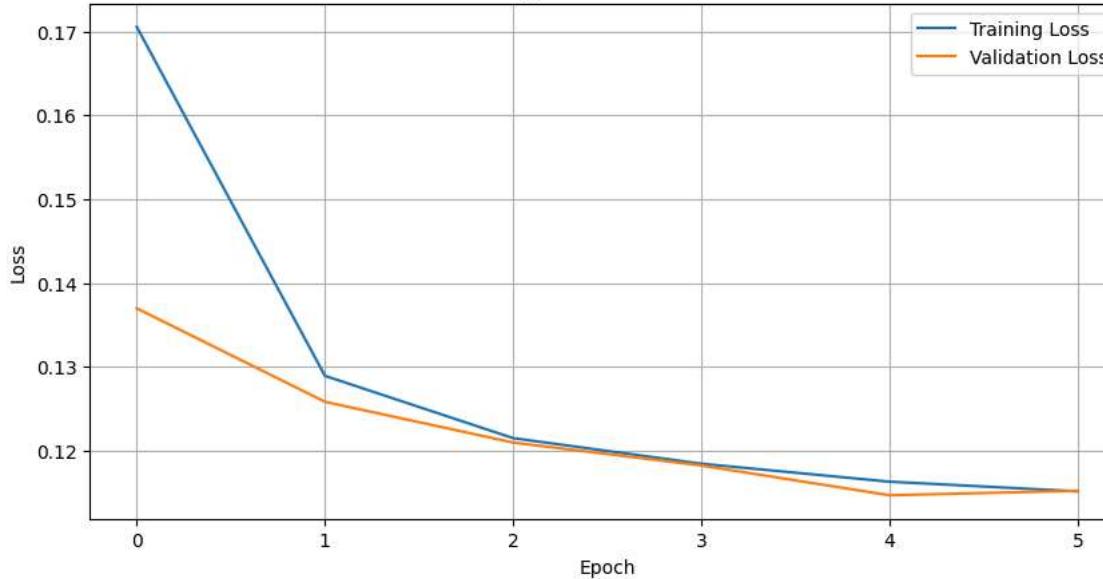
Running validation...

Validation - Epoch 6 average loss: 0.1152

Learning rate: 0.001000

No improvement for 1/10 epochs

Training and Validation Loss



Epoch 7/30

```
-----  
Step 0/750, Loss: 0.1070  
Step 100/750, Loss: 0.1224  
Step 200/750, Loss: 0.1240  
Step 300/750, Loss: 0.1040
```

4/6/25, 10:01 AM

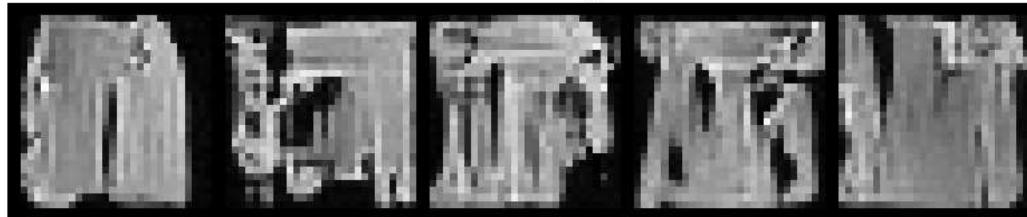
MD_Notebook_OlugbengaAdegoroye_ITAI_2376.ipynb - Colab

Step 400/750, Loss: 0.1258

Step 500/750, Loss: 0.1026

Generating samples...

Generated Samples



Step 600/750, Loss: 0.1180

Step 700/750, Loss: 0.1272

Training - Epoch 7 average loss: 0.1135

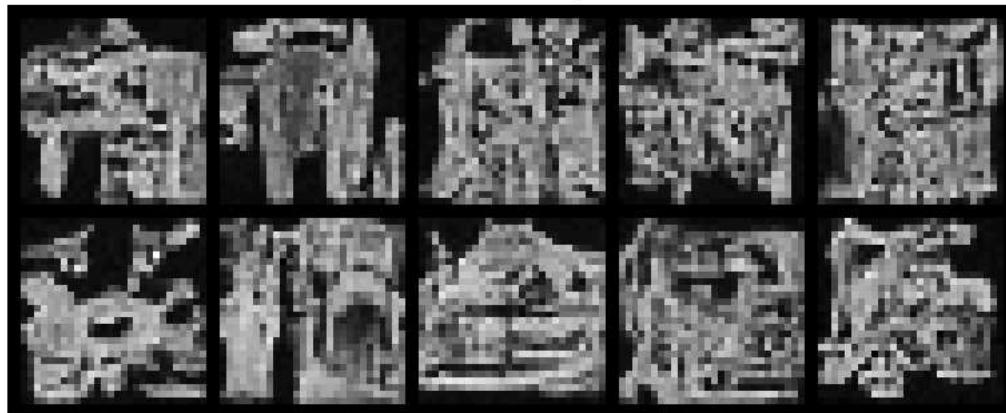
Running validation...

Validation - Epoch 7 average loss: 0.1133

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



Created backup at best_diffusion_model.pt.backup

Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.1133)

Epoch 8/30

Step 0/750, Loss: 0.0875

Step 100/750, Loss: 0.1153

Step 200/750, Loss: 0.1084

Step 300/750, Loss: 0.1258

Step 400/750, Loss: 0.1171

Step 500/750, Loss: 0.1140

Generating samples...

Generated Samples





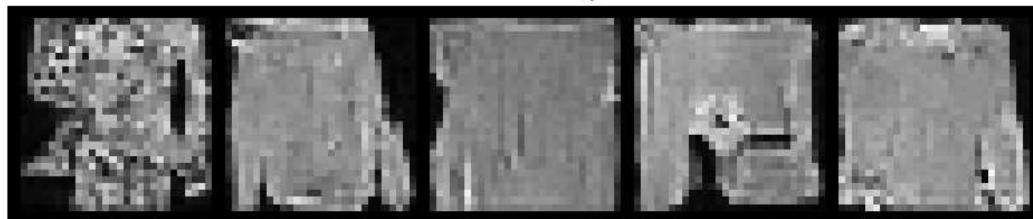
Step 600/750, Loss: 0.1303
Step 700/750, Loss: 0.1125

Training - Epoch 8 average loss: 0.1123
Running validation...
Validation - Epoch 8 average loss: 0.1139
Learning rate: 0.001000
No improvement for 1/10 epochs

Epoch 9/30

Step 0/750, Loss: 0.0855
Step 100/750, Loss: 0.1219
Step 200/750, Loss: 0.0977
Step 300/750, Loss: 0.1172
Step 400/750, Loss: 0.1244
Step 500/750, Loss: 0.0960
Generating samples...

Generated Samples

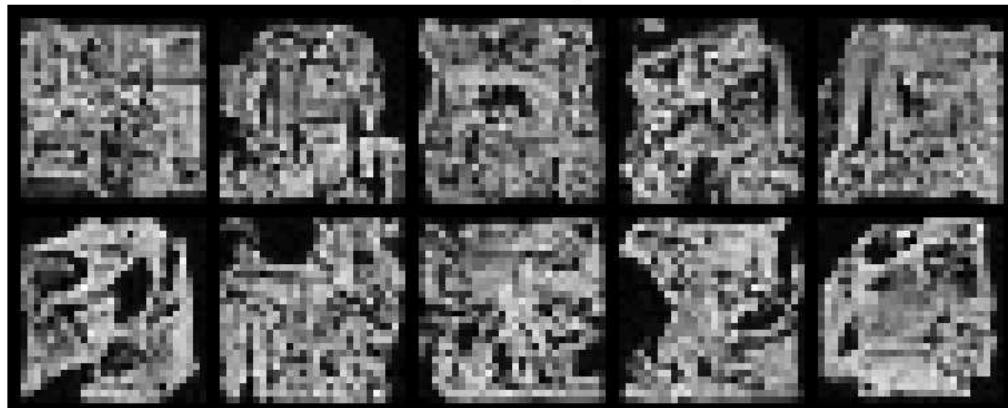


Step 600/750, Loss: 0.1031
Step 700/750, Loss: 0.0984

Training - Epoch 9 average loss: 0.1119
Running validation...
Validation - Epoch 9 average loss: 0.1109
Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



4/6/25, 10:01 AM

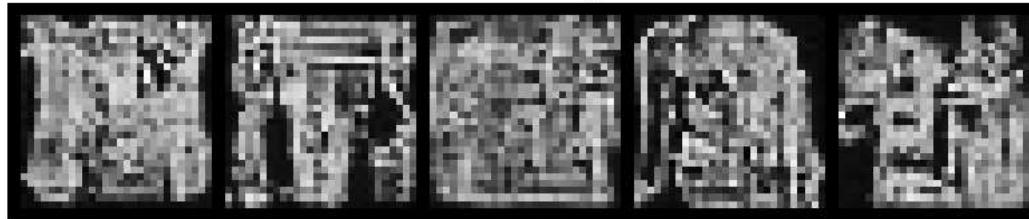
MD_Notebook_OlugbengaAdegoroye_ITAI_2376.ipynb - Colab

```
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1109)
```

Epoch 10/30

```
Step 0/750, Loss: 0.1165
Step 100/750, Loss: 0.0990
Step 200/750, Loss: 0.1133
Step 300/750, Loss: 0.1197
Step 400/750, Loss: 0.0993
Step 500/750, Loss: 0.1185
Generating samples...
```

Generated Samples



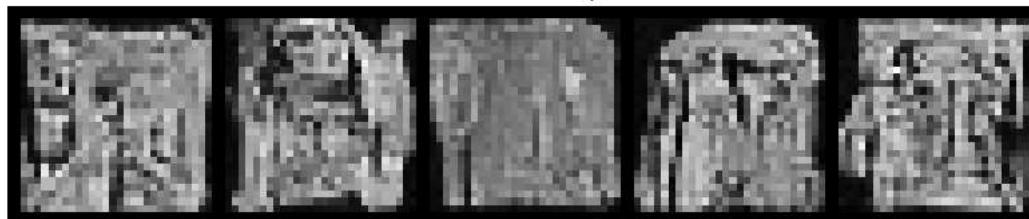
```
Step 600/750, Loss: 0.1022
Step 700/750, Loss: 0.1381
```

```
Training - Epoch 10 average loss: 0.1112
Running validation...
Validation - Epoch 10 average loss: 0.1119
Learning rate: 0.001000
No improvement for 1/10 epochs
```

Epoch 11/30

```
Step 0/750, Loss: 0.1330
Step 100/750, Loss: 0.1299
Step 200/750, Loss: 0.1082
Step 300/750, Loss: 0.1082
Step 400/750, Loss: 0.0997
Step 500/750, Loss: 0.1128
Generating samples...
```

Generated Samples

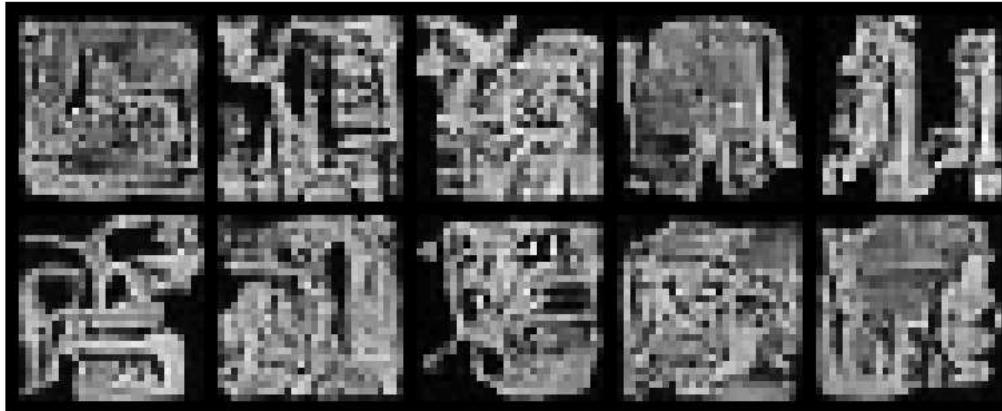


```
Step 600/750, Loss: 0.1457
Step 700/750, Loss: 0.1291
```

```
Training - Epoch 11 average loss: 0.1105
Running validation...
Validation - Epoch 11 average loss: 0.1107
Learning rate: 0.001000
```

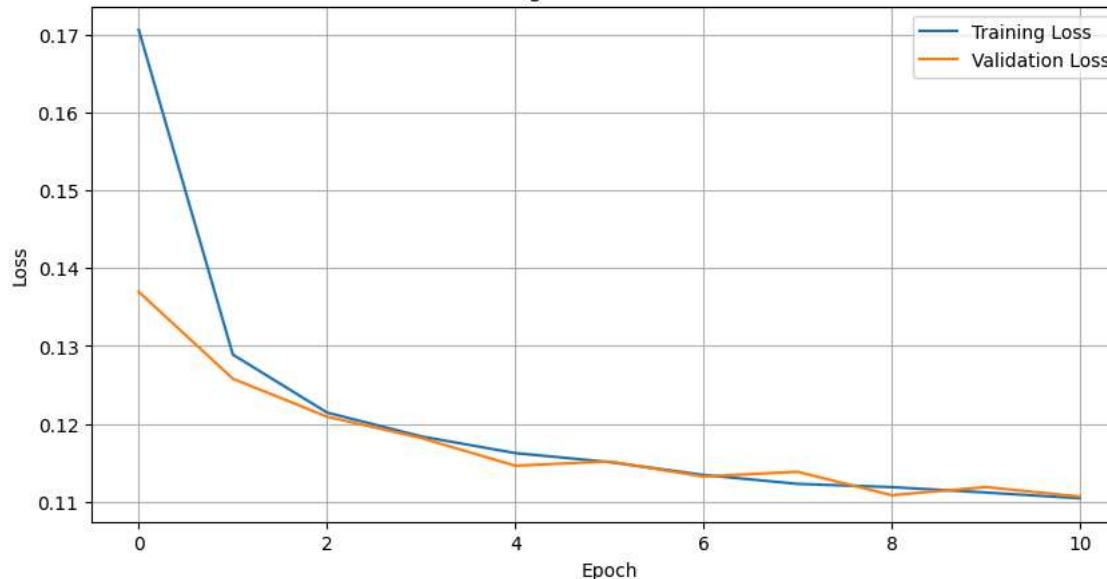
Generating samples for visual progress check...

Generated Samples



Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1107)

Training and Validation Loss

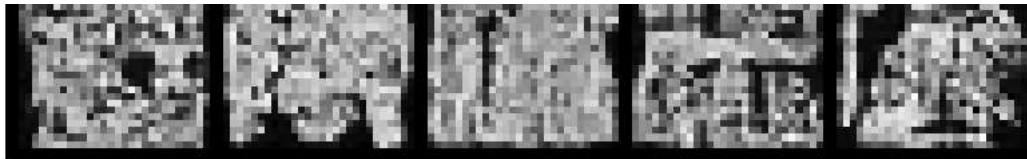


Epoch 12/30

Step 0/750, Loss: 0.1028
Step 100/750, Loss: 0.1108
Step 200/750, Loss: 0.1170
Step 300/750, Loss: 0.1295
Step 400/750, Loss: 0.1137
Step 500/750, Loss: 0.1006
Generating samples...

Generated Samples





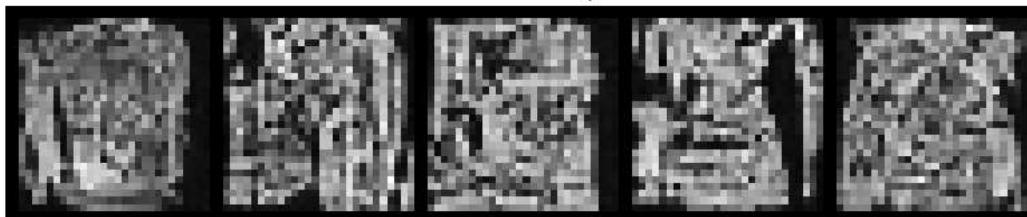
Step 600/750, Loss: 0.0844
Step 700/750, Loss: 0.1296

```
Training - Epoch 12 average loss: 0.1105
Running validation...
Validation - Epoch 12 average loss: 0.1104
Learning rate: 0.001000
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1104)
```

Epoch 13/30

```
-----  
Step 0/750, Loss: 0.1211  
Step 100/750, Loss: 0.1453  
Step 200/750, Loss: 0.0918  
Step 300/750, Loss: 0.1426  
Step 400/750, Loss: 0.1175  
Step 500/750, Loss: 0.0957  
Generating samples...
```

Generated Samples

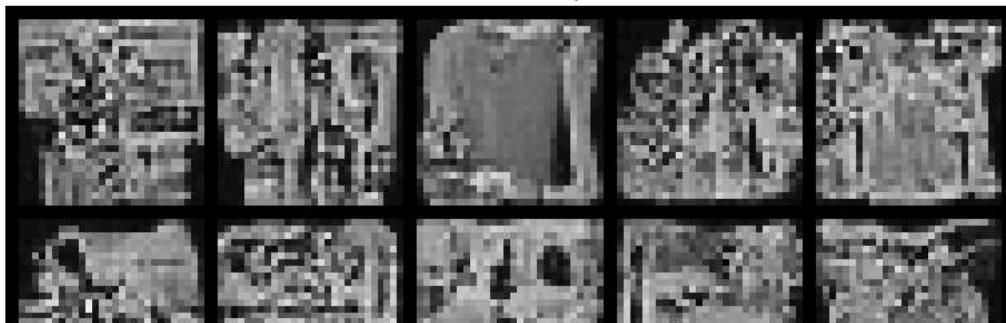


Step 600/750, Loss: 0.0817
Step 700/750, Loss: 0.1205

```
Training - Epoch 13 average loss: 0.1105
Running validation...
Validation - Epoch 13 average loss: 0.1108
Learning rate: 0.001000
```

Generating samples for visual progress check...

Generated Samples



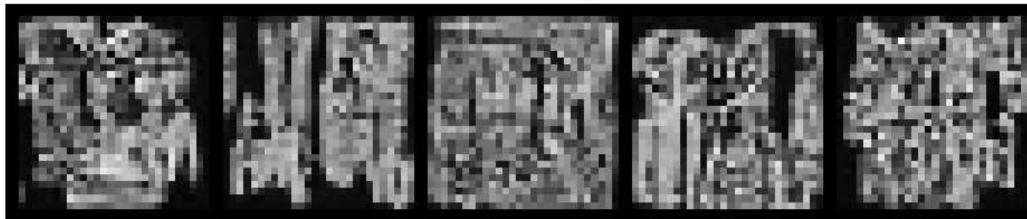


No improvement for 1/10 epochs

Epoch 14/30

```
-----  
Step 0/750, Loss: 0.0824  
Step 100/750, Loss: 0.1252  
Step 200/750, Loss: 0.0904  
Step 300/750, Loss: 0.0626  
Step 400/750, Loss: 0.1082  
Step 500/750, Loss: 0.1208  
Generating samples...
```

Generated Samples



Step 600/750, Loss: 0.1254
Step 700/750, Loss: 0.0954

Training - Epoch 14 average loss: 0.1095

Running validation...

Validation - Epoch 14 average loss: 0.1109

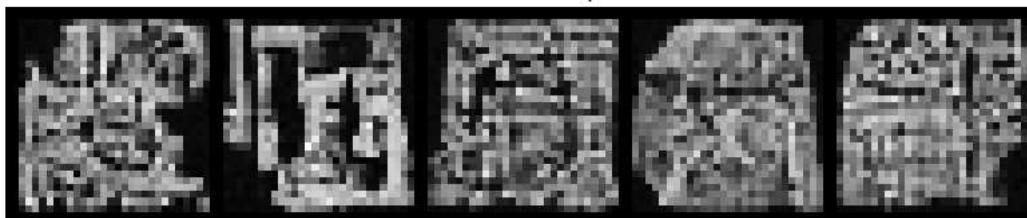
Learning rate: 0.001000

No improvement for 2/10 epochs

Epoch 15/30

```
-----  
Step 0/750, Loss: 0.1403  
Step 100/750, Loss: 0.1270  
Step 200/750, Loss: 0.0945  
Step 300/750, Loss: 0.1046  
Step 400/750, Loss: 0.0872  
Step 500/750, Loss: 0.1023  
Generating samples...
```

Generated Samples



Step 600/750, Loss: 0.1368
Step 700/750, Loss: 0.1041

Training - Epoch 15 average loss: 0.1089

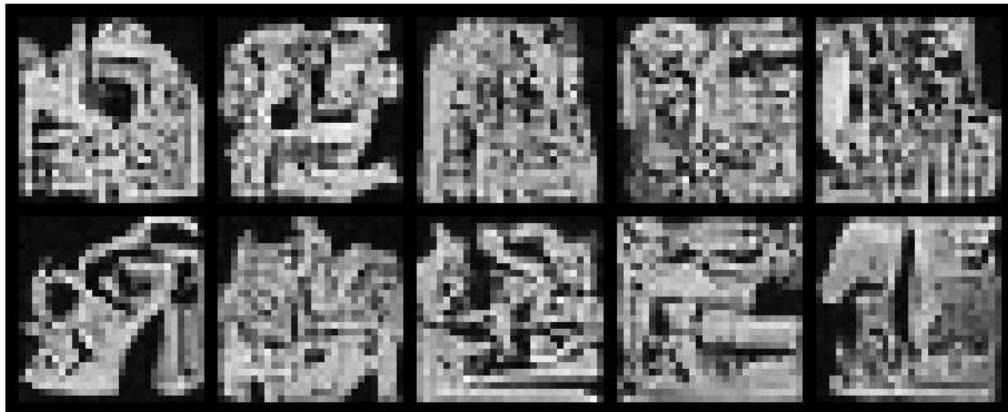
Running validation...

Validation - Epoch 15 average loss: 0.1079

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



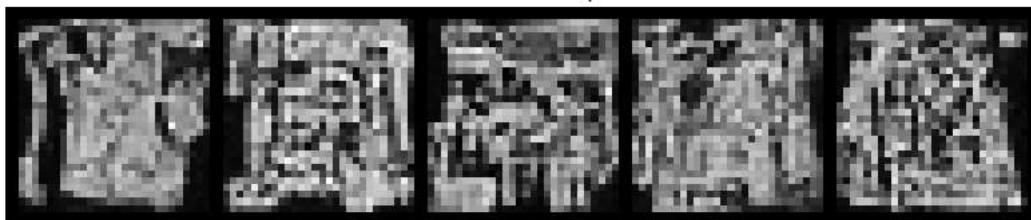
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1079)

Epoch 16/30

```
-----  
Step 0/750, Loss: 0.1000  
Step 100/750, Loss: 0.1002  
Step 200/750, Loss: 0.0929  
Step 300/750, Loss: 0.0921  
Step 400/750, Loss: 0.1171  
Step 500/750, Loss: 0.1149
```

Generating samples...

Generated Samples



Step 600/750, Loss: 0.0894
Step 700/750, Loss: 0.1119

Training - Epoch 16 average loss: 0.1088

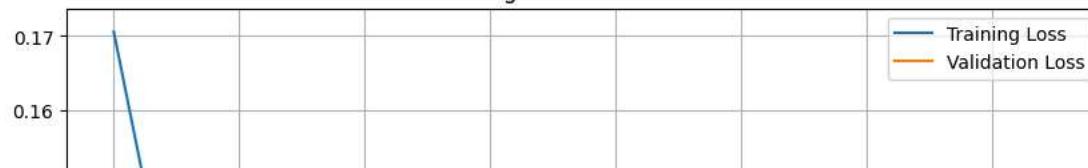
Running validation...

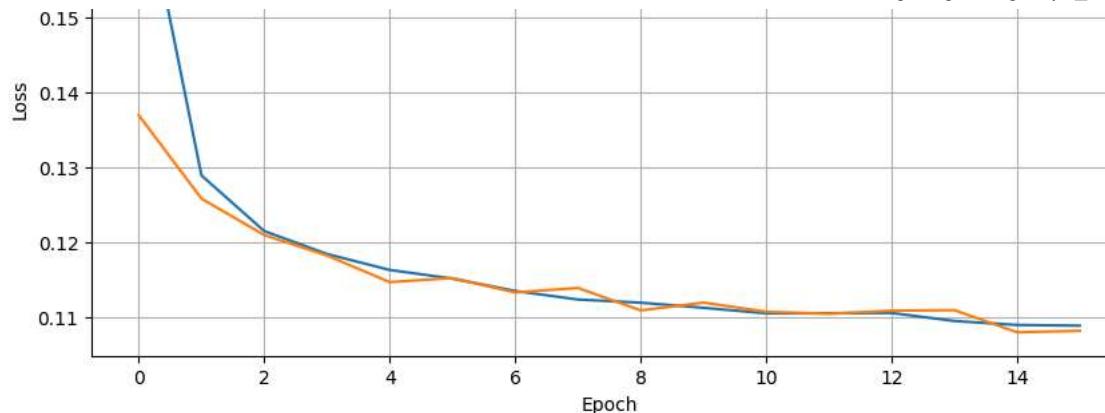
Validation - Epoch 16 average loss: 0.1081

Learning rate: 0.001000

No improvement for 1/10 epochs

Training and Validation Loss

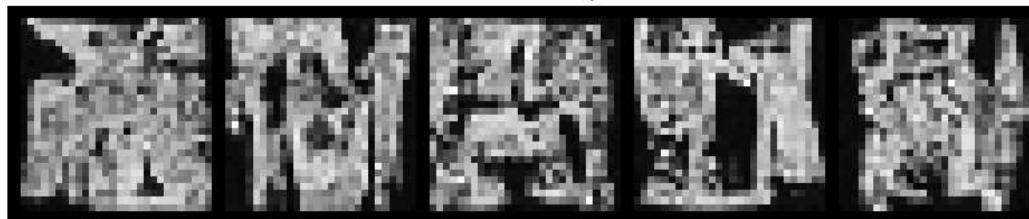




Epoch 17/30

Step 0/750, Loss: 0.0853
Step 100/750, Loss: 0.1196
Step 200/750, Loss: 0.1080
Step 300/750, Loss: 0.0800
Step 400/750, Loss: 0.1109
Step 500/750, Loss: 0.1252
Generating samples...

Generated Samples



Step 600/750, Loss: 0.1174
Step 700/750, Loss: 0.0941

Training - Epoch 17 average loss: 0.1097

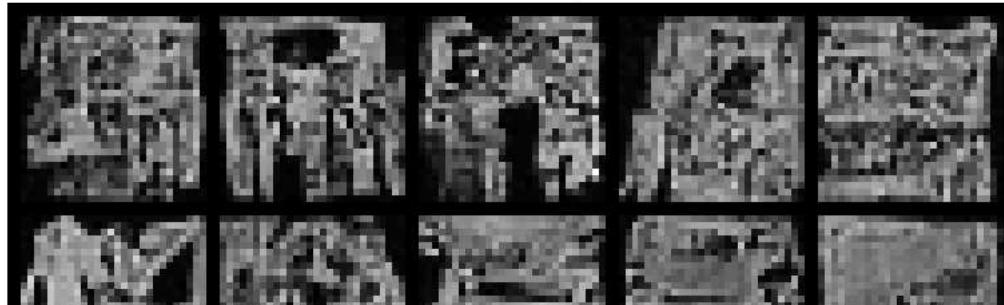
Running validation...

Validation - Epoch 17 average loss: 0.1077

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



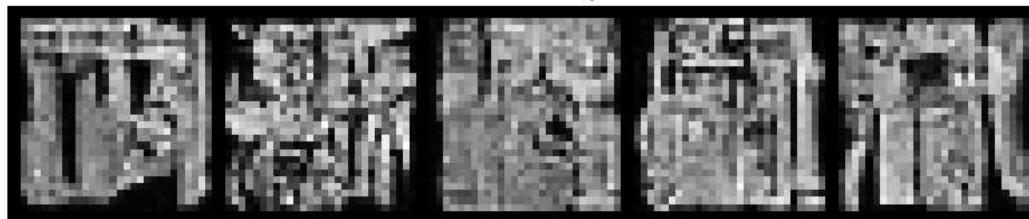


```
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1077)
```

Epoch 18/30

```
-----  
Step 0/750, Loss: 0.1068  
Step 100/750, Loss: 0.1069  
Step 200/750, Loss: 0.1147  
Step 300/750, Loss: 0.1375  
Step 400/750, Loss: 0.0949  
Step 500/750, Loss: 0.0956  
Generating samples...
```

Generated Samples



```
Step 600/750, Loss: 0.0997  
Step 700/750, Loss: 0.1111
```

Training - Epoch 18 average loss: 0.1071

Running validation...

Validation - Epoch 18 average loss: 0.1102

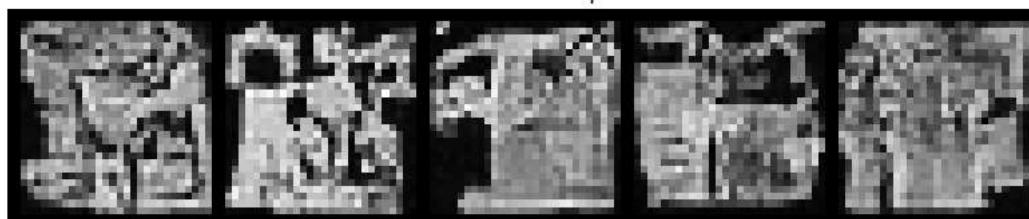
Learning rate: 0.001000

No improvement for 1/10 epochs

Epoch 19/30

```
-----  
Step 0/750, Loss: 0.1066  
Step 100/750, Loss: 0.1367  
Step 200/750, Loss: 0.1095  
Step 300/750, Loss: 0.1117  
Step 400/750, Loss: 0.1083  
Step 500/750, Loss: 0.1032  
Generating samples...
```

Generated Samples



```
Step 600/750, Loss: 0.1174  
Step 700/750, Loss: 0.0958
```

Training - Epoch 19 average loss: 0.1080

Running validation...

4/6/25, 10:01 AM

MD_Notebook_OlugbengaAdegoroye_ITAI_2376.ipynb - Colab

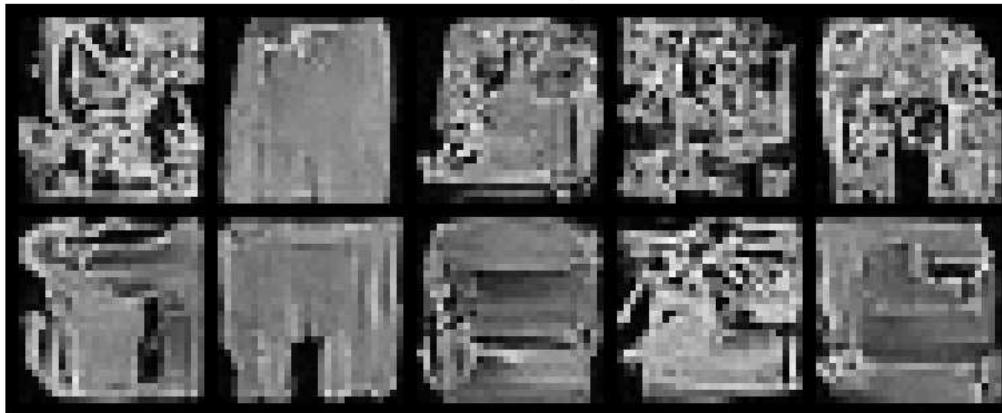
Running validation...

Validation - Epoch 19 average loss: 0.1063

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples

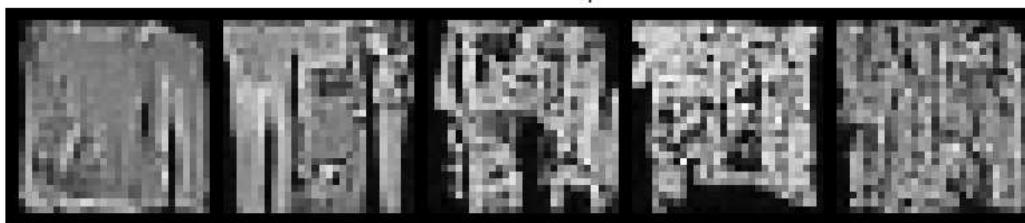


Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1063)

Epoch 20/30

Step 0/750, Loss: 0.1212
Step 100/750, Loss: 0.1424
Step 200/750, Loss: 0.0863
Step 300/750, Loss: 0.1020
Step 400/750, Loss: 0.1079
Step 500/750, Loss: 0.1030
Generating samples...

Generated Samples



Step 600/750, Loss: 0.1008
Step 700/750, Loss: 0.0872

Training - Epoch 20 average loss: 0.1078

Running validation...

Validation - Epoch 20 average loss: 0.1082

Learning rate: 0.001000

No improvement for 1/10 epochs

Epoch 21/30

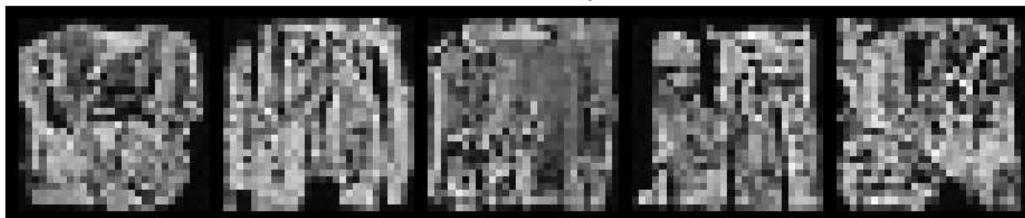
Step 0/750, Loss: 0.1134
Step 100/750, Loss: 0.0854
Step 200/750, Loss: 0.0961
Step 300/750, Loss: 0.1021

4/6/25, 10:01 AM

MD_Notebook_OlugbengaAdegoroye_ITAI_2376.ipynb - Colab

```
Step 200/750, Loss: 0.1054
Step 400/750, Loss: 0.1034
Step 500/750, Loss: 0.1191
Generating samples...
```

Generated Samples



```
Step 600/750, Loss: 0.0934
Step 700/750, Loss: 0.1113
```

Training - Epoch 21 average loss: 0.1078

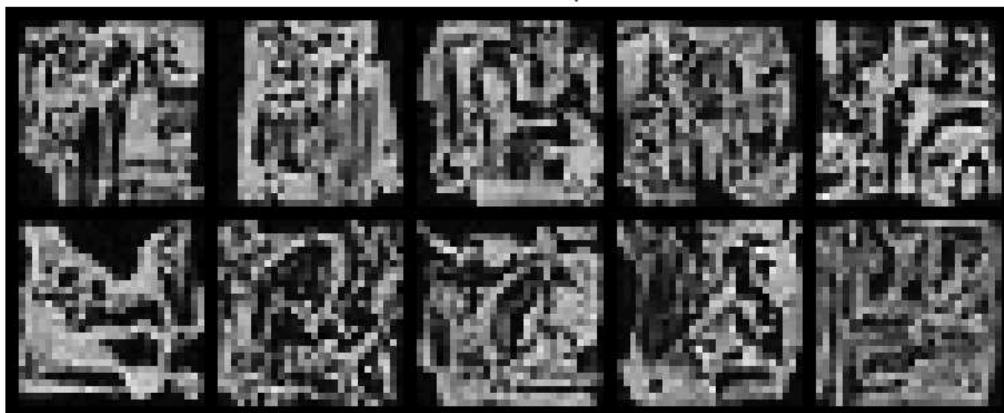
Running validation...

Validation - Epoch 21 average loss: 0.1076

Learning rate: 0.001000

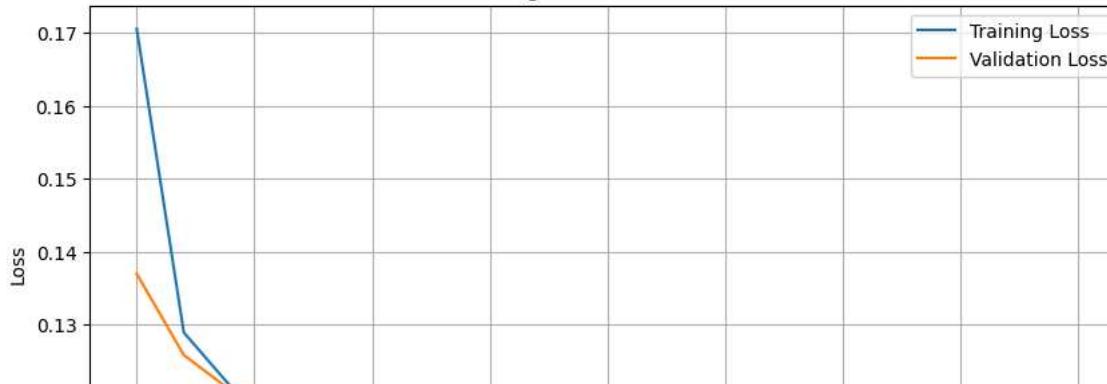
Generating samples for visual progress check...

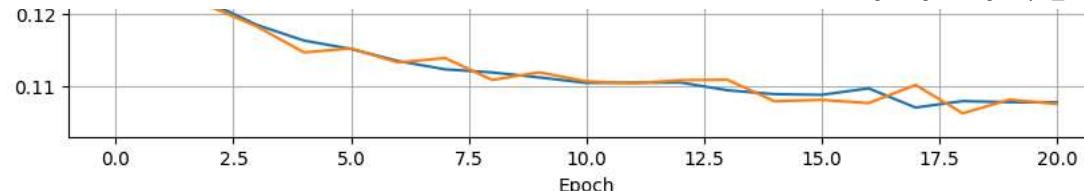
Generated Samples



No improvement for 2/10 epochs

Training and Validation Loss

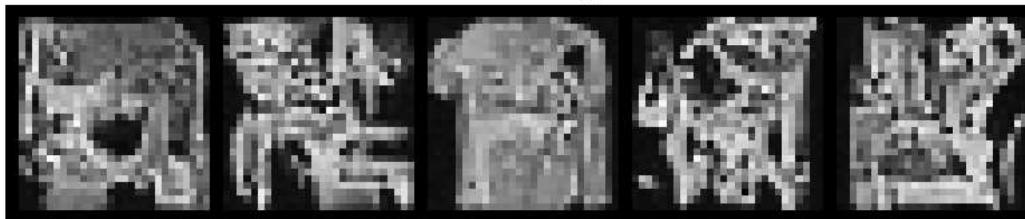




Epoch 22/30

```
Step 0/750, Loss: 0.1070
Step 100/750, Loss: 0.0957
Step 200/750, Loss: 0.1099
Step 300/750, Loss: 0.0993
Step 400/750, Loss: 0.1129
Step 500/750, Loss: 0.0990
Generating samples...
```

Generated Samples



```
Step 600/750, Loss: 0.1042
Step 700/750, Loss: 0.1216
```

Training - Epoch 22 average loss: 0.1070

Running validation...

Validation - Epoch 22 average loss: 0.1102

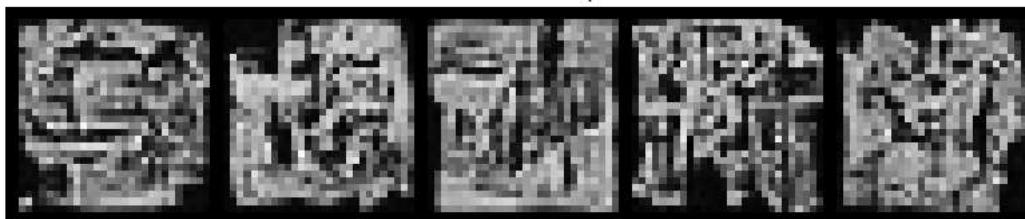
Learning rate: 0.001000

No improvement for 3/10 epochs

Epoch 23/30

```
Step 0/750, Loss: 0.1072
Step 100/750, Loss: 0.0842
Step 200/750, Loss: 0.0933
Step 300/750, Loss: 0.1221
Step 400/750, Loss: 0.1136
Step 500/750, Loss: 0.1169
Generating samples...
```

Generated Samples



```
Step 600/750, Loss: 0.1440
Step 700/750, Loss: 0.0957
```

Training - Epoch 23 average loss: 0.1076

Running validation...
Validation - Epoch 23 average loss: 0.1090
Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples

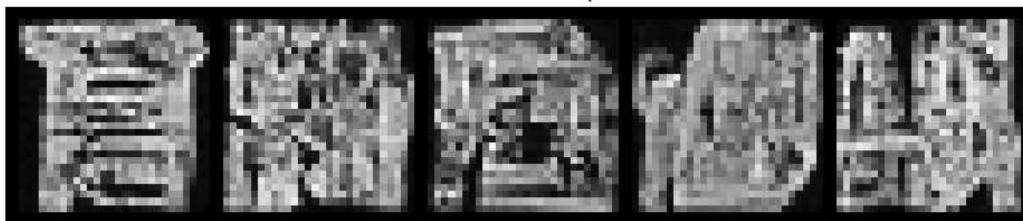


No improvement for 4/10 epochs

Epoch 24/30

Step 0/750, Loss: 0.1168
Step 100/750, Loss: 0.0971
Step 200/750, Loss: 0.1281
Step 300/750, Loss: 0.1010
Step 400/750, Loss: 0.1071
Step 500/750, Loss: 0.1236
Generating samples...

Generated Samples



Step 600/750, Loss: 0.0942
Step 700/750, Loss: 0.0943

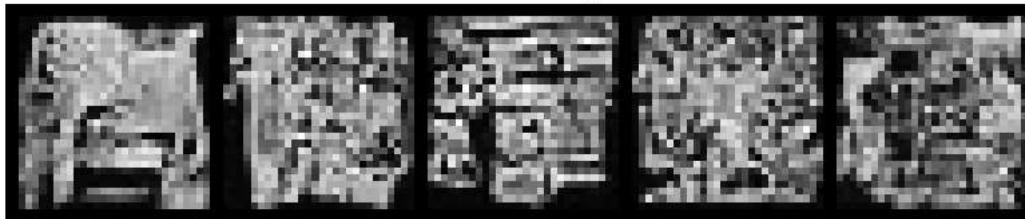
Training - Epoch 24 average loss: 0.1071
Running validation...
Validation - Epoch 24 average loss: 0.1084
Learning rate: 0.001000
No improvement for 5/10 epochs

Epoch 25/30

Step 0/750, Loss: 0.0869
Step 100/750, Loss: 0.1171
Step 200/750, Loss: 0.0967
Step 300/750, Loss: 0.1243
Step 400/750, Loss: 0.1108

Step 500/750, Loss: 0.1003
Generating samples...

Generated Samples

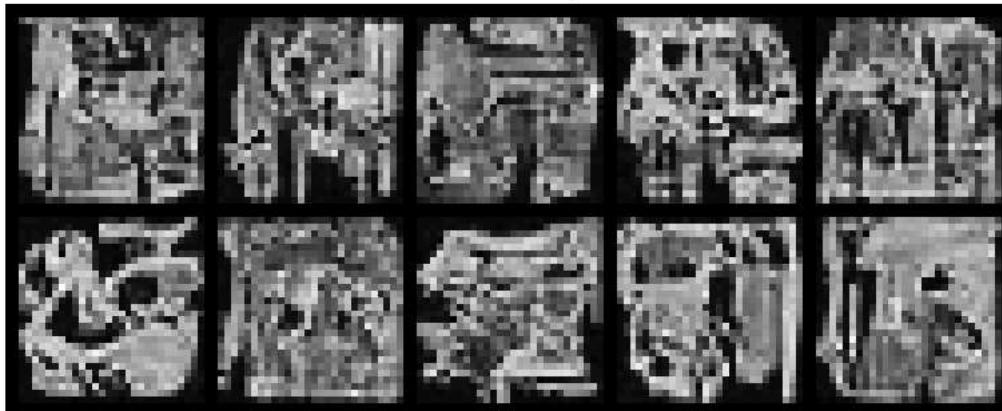


Step 600/750, Loss: 0.0874
Step 700/750, Loss: 0.1026

Training - Epoch 25 average loss: 0.1066
Running validation...
Validation - Epoch 25 average loss: 0.1073
Learning rate: 0.000500

Generating samples for visual progress check...

Generated Samples

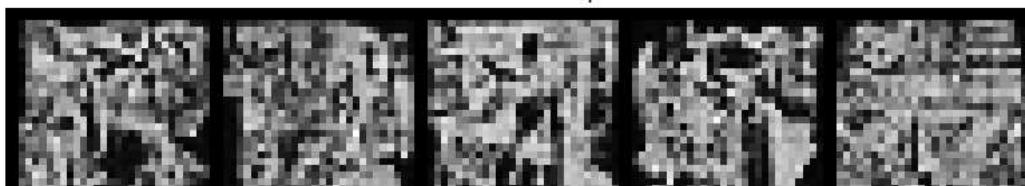


No improvement for 6/10 epochs

Epoch 26/30

Step 0/750, Loss: 0.1053
Step 100/750, Loss: 0.0661
Step 200/750, Loss: 0.1191
Step 300/750, Loss: 0.0974
Step 400/750, Loss: 0.1166
Step 500/750, Loss: 0.1135
Generating samples...

Generated Samples



Step 600/750, Loss: 0.1006
Step 700/750, Loss: 0.1265

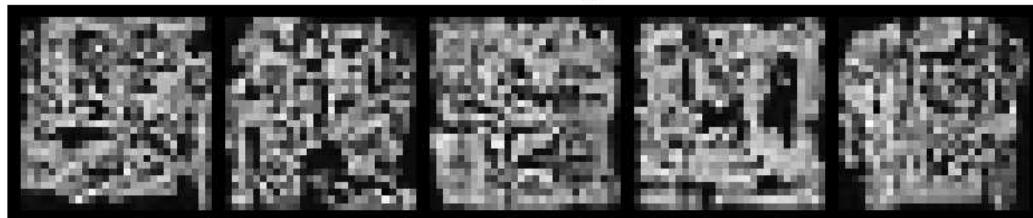
Training - Epoch 26 average loss: 0.1041
Running validation...
Validation - Epoch 26 average loss: 0.1038
Learning rate: 0.000500
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1038)



Epoch 27/30

Step 0/750, Loss: 0.0972
Step 100/750, Loss: 0.1021
Step 200/750, Loss: 0.1205
Step 300/750, Loss: 0.0987
Step 400/750, Loss: 0.0860
Step 500/750, Loss: 0.1012
Generating samples...

Generated Samples



Step 600/750, Loss: 0.1149
Step 700/750, Loss: 0.1040

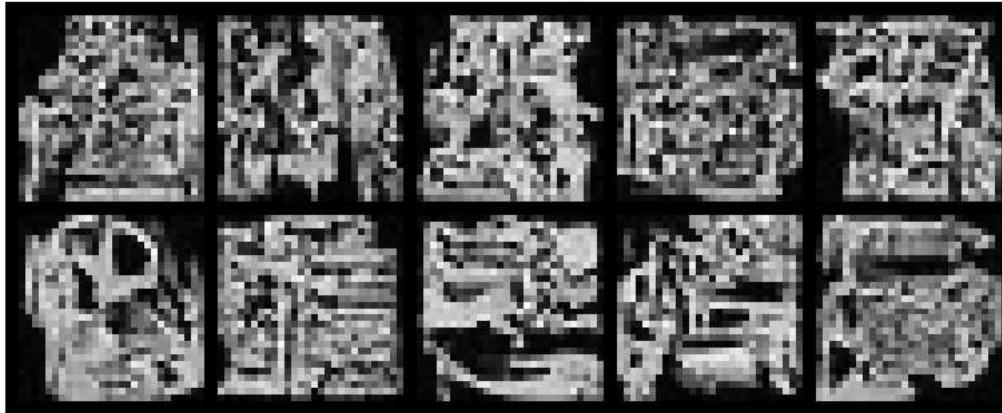
4/6/25, 10:01 AM

MD_Notebook_OlugbengaAdegoroye_ITAI_2376.ipynb - Colab

```
Training - Epoch 27 average loss: 0.1047
Running validation...
Validation - Epoch 27 average loss: 0.1049
Learning rate: 0.000500
```

Generating samples for visual progress check...

Generated Samples

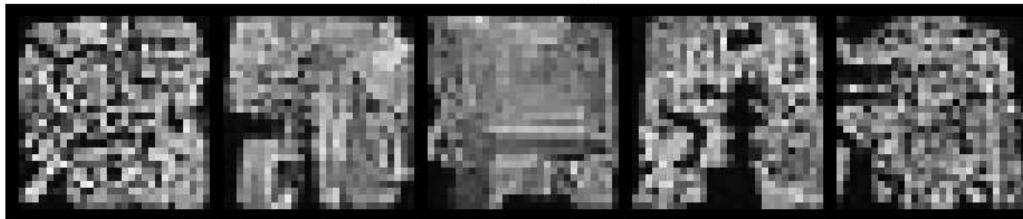


No improvement for 1/10 epochs

Epoch 28/30

```
Step 0/750, Loss: 0.1035
Step 100/750, Loss: 0.1027
Step 200/750, Loss: 0.1051
Step 300/750, Loss: 0.1289
Step 400/750, Loss: 0.1156
Step 500/750, Loss: 0.1082
Generating samples...
```

Generated Samples



```
Step 600/750, Loss: 0.1219
Step 700/750, Loss: 0.0763
```

```
Training - Epoch 28 average loss: 0.1041
Running validation...
Validation - Epoch 28 average loss: 0.1030
Learning rate: 0.000500
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.1030)
```

Epoch 29/30

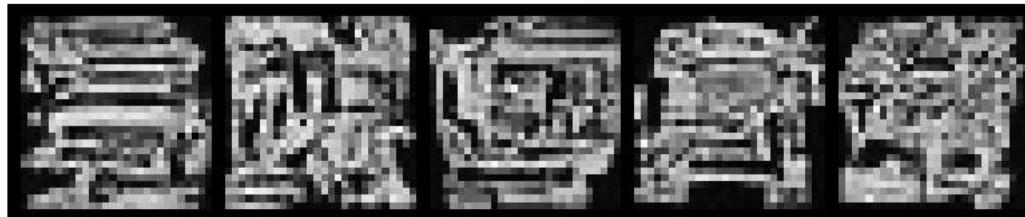
```
Step 0/750, Loss: 0.0909
Step 100/750, Loss: 0.0973
```

4/6/25, 10:01 AM

MD_Notebook_OlugbengaAdegoroye_ITAI_2376.ipynb - Colab

```
Step 200/750, Loss: 0.1133
Step 300/750, Loss: 0.0986
Step 400/750, Loss: 0.1064
Step 500/750, Loss: 0.1047
Generating samples...
```

Generated Samples

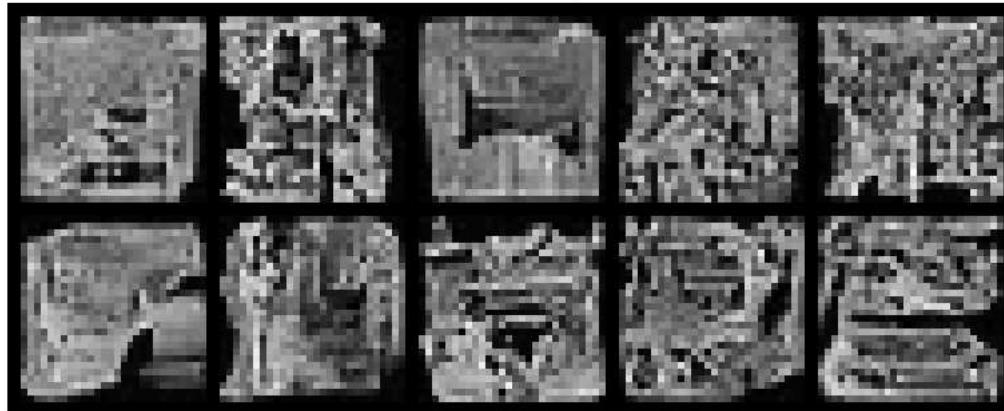


```
Step 600/750, Loss: 0.0989
Step 700/750, Loss: 0.1110
```

```
Training - Epoch 29 average loss: 0.1042
Running validation...
Validation - Epoch 29 average loss: 0.1031
Learning rate: 0.000500
```

Generating samples for visual progress check...

Generated Samples



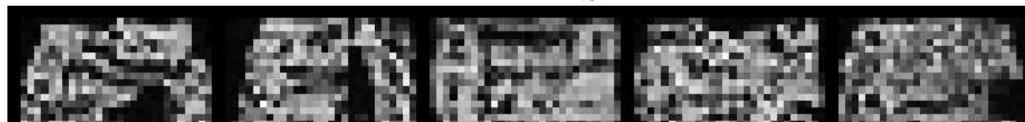
No improvement for 1/10 epochs

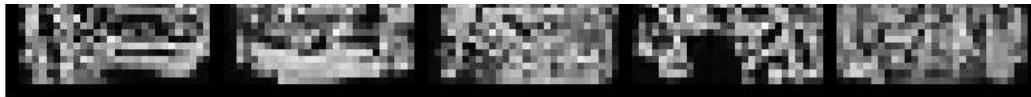
Epoch 30/30

```
-----
```

```
Step 0/750, Loss: 0.1032
Step 100/750, Loss: 0.1086
Step 200/750, Loss: 0.0936
Step 300/750, Loss: 0.1062
Step 400/750, Loss: 0.1053
Step 500/750, Loss: 0.0999
Generating samples...
```

Generated Samples





Step 600/750, Loss: 0.0919
Step 700/750, Loss: 0.0991

Training - Epoch 30 average loss: 0.1045

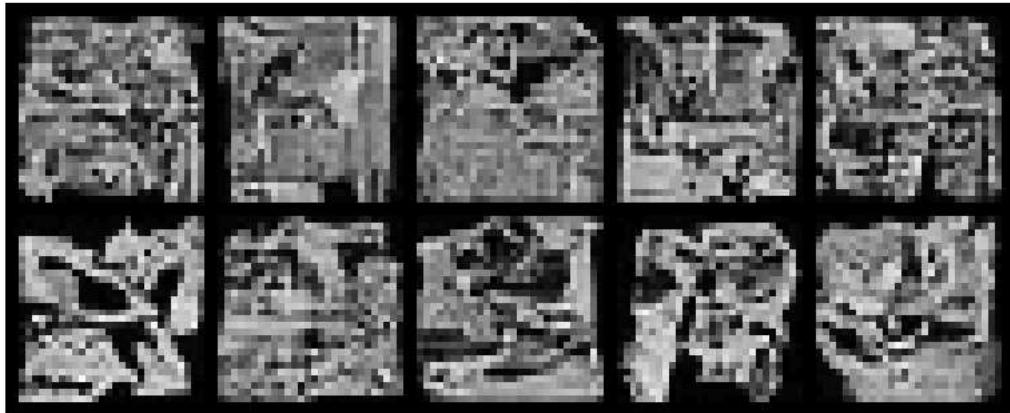
Running validation...

Validation - Epoch 30 average loss: 0.1053

Learning rate: 0.000500

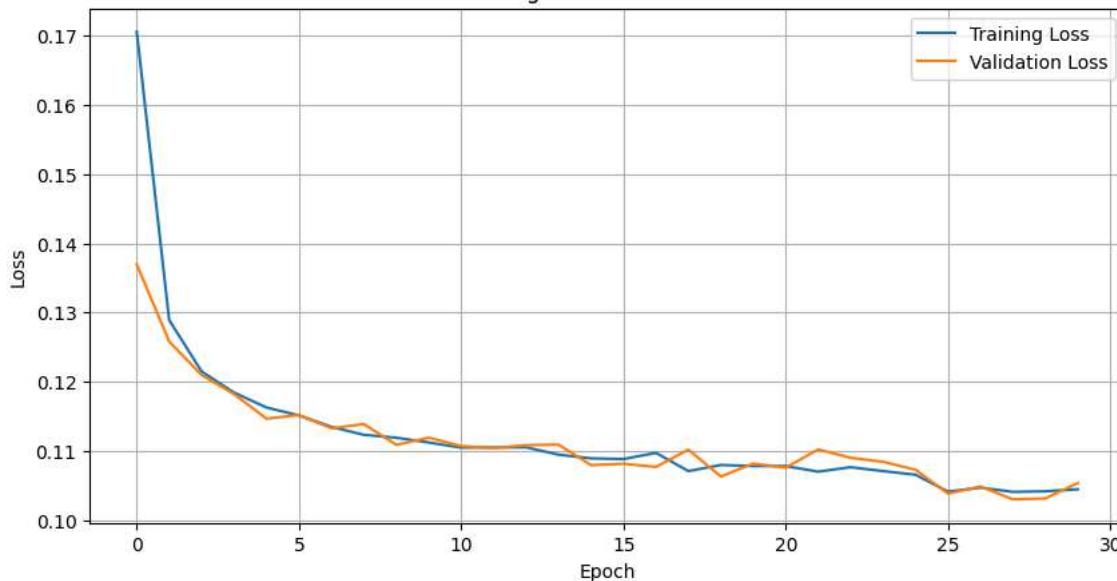
Generating samples for visual progress check...

Generated Samples



No improvement for 2/10 epochs

Training and Validation Loss



=====

4/6/25, 10:01 AM

MD_Notebook_OlugbengaAdegoroye_ITAI_2376.ipynb - Colab

TRAINING COMPLETE

=====

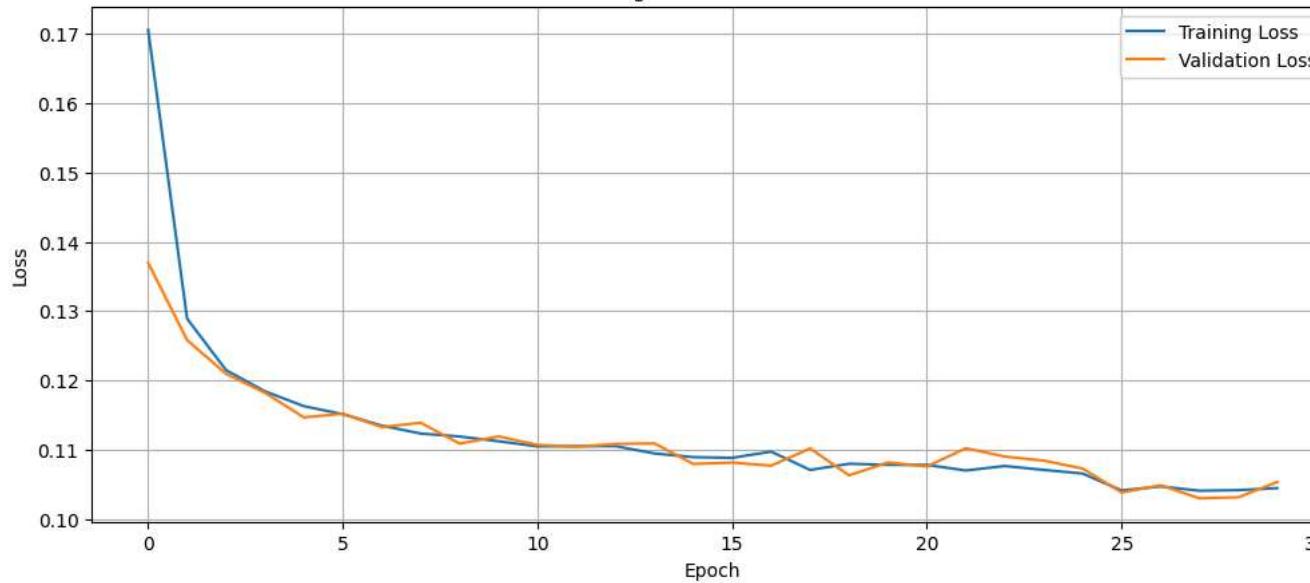
Best validation loss: 0.1030

Generating final samples...

Generated Samples



Training and Validation Loss

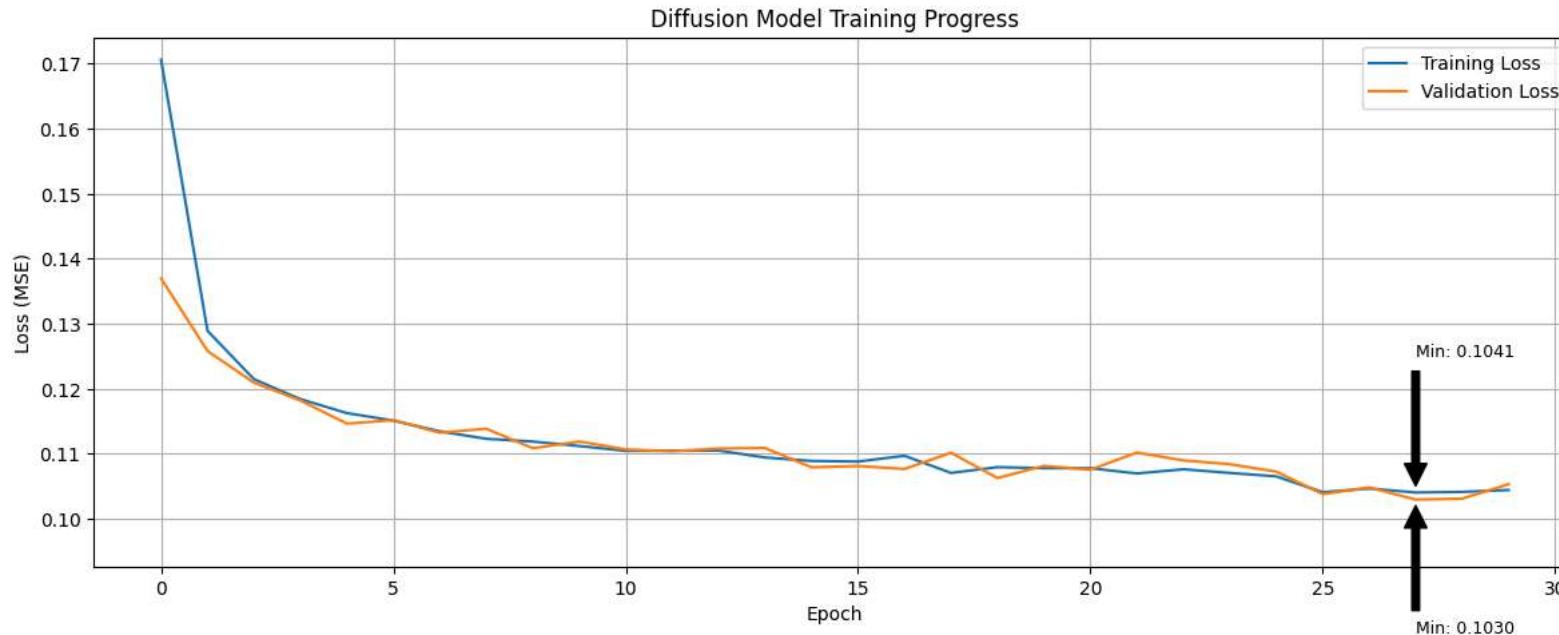


```

1 # Plot training progress
2 plt.figure(figsize=(12, 5))
3
4 # Plot training and validation losses for comparison
5 plt.plot(train_losses, label='Training Loss')
6 if len(val_losses) > 0: # Only plot validation if it exists
7     plt.plot(val_losses, label='Validation Loss')
8
9 # Improve the plot with better labels and styling
10 plt.title('Diffusion Model Training Progress')
11 plt.xlabel('Epoch')
12 plt.ylabel('Loss (MSE)')
13 plt.legend()
14 plt.grid(True)
15
16 # Add annotations for key points
17 if len(train_losses) > 1:
18     min_train_idx = train_losses.index(min(train_losses))
19     plt.annotate(f'Min: {min(train_losses):.4f}',
20                  xy=(min_train_idx, min(train_losses)),
21                  xytext=(min_train_idx, min(train_losses)*1.2),
22                  arrowprops=dict(facecolor='black', shrink=0.05),
23                  fontsize=9)
24
25 # Add validation min point if available
26 if len(val_losses) > 1:
27     min_val_idx = val_losses.index(min(val_losses))
28     plt.annotate(f'Min: {min(val_losses):.4f}',
29                  xy=(min_val_idx, min(val_losses)),
30                  xytext=(min_val_idx, min(val_losses)*0.8),
31                  arrowprops=dict(facecolor='black', shrink=0.05),
32                  fontsize=9)
33
34 # Set y-axis to start from 0 or slightly lower than min value
35 plt.ylim(bottom=max(0, min(min(train_losses) if train_losses else float('inf'),
36                           min(val_losses) if val_losses else float('inf'))*0.9)))
37
38 plt.tight_layout()
39 plt.show()
40
41 # Add statistics summary for students to analyze
42 print("\nTraining Statistics:")
43 print("-" * 30)
44 if train_losses:
45     print(f"Starting training loss: {train_losses[0]:.4f}")
46     print(f"Final training loss: {train_losses[-1]:.4f}")
47     print(f"Best training loss: {min(train_losses):.4f}")
48     print(f"Training loss improvement: {((train_losses[0] - min(train_losses)) / train_losses[0]) * 100:.1f}%")
49
50 if val_losses:
51     print("\nValidation Statistics:")
52     print("-" * 30)
53     print(f"Starting validation loss: {val_losses[0]:.4f}")
54     print(f"Final validation loss: {val_losses[-1]:.4f}")
55     print(f"Best validation loss: {min(val_losses):.4f}")
56
57 # STUDENT EXERCISE:

```

```
58 # 1. Try modifying this plot to show a smoothed version of the losses
59 # 2. Create a second plot showing the ratio of validation to training loss
60 #     (which can indicate overfitting when the ratio increases)
```



Training Statistics:

```
-----  
Starting training loss: 0.1706  
Final training loss: 0.1045  
Best training loss: 0.1041  
Training loss improvement: 39.0%
```

Validation Statistics:

```
-----  
Starting validation loss: 0.1370  
Final validation loss: 0.1053  
Best validation loss: 0.1030
```

```
1 # STUDENT EXERCISE:  
2 # 1. Try modifying this plot to show a smoothed version of the losses  
3 # 2. Create a second plot showing the ratio of validation to training loss  
4 #     (which can indicate overfitting when the ratio increases)  
5  
6 print("\nStudent Exercise: Enhancing Loss Visualization")  
7  
8 # Using pre-defined libraries from Step 1  
9 import torch  
10 import matplotlib.pyplot as plt  
11 import numpy as np  
12  
13 # Using pre-defined EPOCHS from Step 2 (Fashion-MNIST)  
14 epochs = EPOCHS # 30, as set in the Fashion-MNIST section  
15
```

```
16 # Check for existing loss data from training
17 try:
18     assert len(train_losses) == epochs, "train_losses length doesn't match EPOCHS"
19     assert len(val_losses) == epochs, "val_losses length doesn't match EPOCHS"
20 except NameError:
21     # Fallback: Simulate losses if not defined (due to truncation)
22     print("Warning: train_losses and val_losses not found. Using simulated data.")
23     print("Replace with actual losses from your training loop.")
24     np.random.seed(42) # Reproducible simulation
25     train_losses = [max(0.5 - 0.015 * i + np.random.normal(0, 0.02), 0.1) for i in range(epochs)]
26     val_losses = [max(0.6 - 0.012 * i + np.random.normal(0, 0.03), 0.15) for i in range(epochs)]
27 except AssertionError as e:
28     print(f"Error: {e}. Ensure losses match EPOCHS ({epochs}). Using available data as-is.")
29
30 # Function to compute moving average for smoothing
31 def moving_average(data, window_size=5):
32     """Compute a moving average of the data with a given window size."""
33     return np.convolve(data, np.ones(window_size) / window_size, mode='valid')
34
35 # Parameters for smoothing
36 window_size = 5 # Adjustable; smaller windows track noise, larger ones smooth more
37 smooth_epochs = range(window_size - 1, epochs) # Adjust x-axis for smoothed data length
38
39 train_losses = np.array(train_losses) if not isinstance(train_losses, np.ndarray) else train_losses
40 val_losses = np.array(val_losses) if not isinstance(val_losses, np.ndarray) else val_losses
41 train_losses_smooth = moving_average(train_losses, window_size)
42 val_losses_smooth = moving_average(val_losses, window_size)
43
44 # Compute validation-to-training loss ratio (using raw data to maintain length)
45 loss_ratio = [val / train if train > 0 else 1.0 for val, train in zip(val_losses, train_losses)]
46
47 # Create subplots
48 plt.figure(figsize=(15, 5))
49
50 # Plot 1: Smoothed Training and Validation Losses
51 plt.subplot(1, 2, 1)
52 plt.plot(range(epochs), train_losses, label="Raw Training Loss", alpha=0.3, color='blue')
53 plt.plot(range(epochs), val_losses, label="Raw Validation Loss", alpha=0.3, color='orange')
54 plt.plot(smooth_epochs, train_losses_smooth, label="Smoothed Training Loss", color='blue')
55 plt.plot(smooth_epochs, val_losses_smooth, label="Smoothed Validation Loss", color='orange')
56 plt.title("Training and Validation Losses (Smoothed)")
57 plt.xlabel("Epoch")
58 plt.ylabel("Loss")
59 plt.legend()
60 plt.grid(True)
61
62 # Plot 2: Validation-to-Training Loss Ratio
63 plt.subplot(1, 2, 2)
64 plt.plot(range(epochs), loss_ratio, label="Val/Train Loss Ratio", color='green')
65 plt.axhline(y=1.0, color='gray', linestyle='--', alpha=0.5)
66 plt.title("Validation to Training Loss Ratio")
67 plt.xlabel("Epoch")
68 plt.ylabel("Ratio")
69 plt.legend()
70 plt.grid(True)
71
72 plt.tight_layout()
73 plt.show()
```

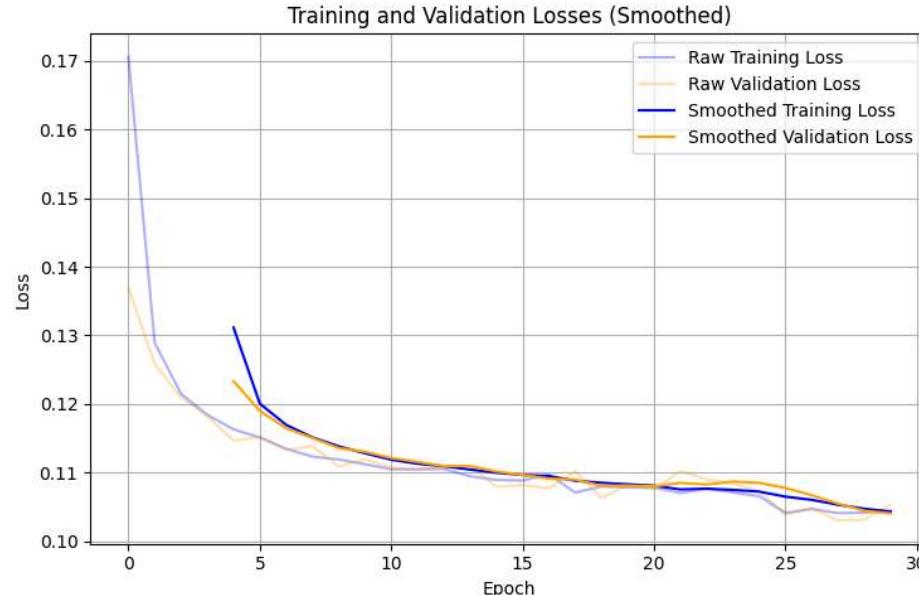
```

75 print("Observations:")
76 print(f"- Smoothed losses use a {window_size}-epoch moving average to reduce noise.")
77 print("- Raw losses are shown faintly for comparison.")
78 print("- Loss Ratio > 1 suggests potential overfitting; a rising trend is a stronger indicator.")
79
80 # Clean up GPU memory, consistent with other snippets
81 if torch.cuda.is_available():
82     torch.cuda.empty_cache()

```



Student Exercise: Enhancing Loss Visualization



Observations:

- Smoothed losses use a 5-epoch moving average to reduce noise.
- Raw losses are shown faintly for comparison.
- Loss Ratio > 1 suggests potential overfitting; a rising trend is a stronger indicator.

▼ Step 6: Generating New Images

Now that our model is trained, let's generate some new images! We can:

1. Generate specific numbers
2. Generate multiple versions of each number
3. See how the generation process works step by step

```

1 def generate_number(model, number, n_samples=4):
2     """
3         Generate multiple versions of a specific number using the diffusion model.
4
5     Args:

```

```

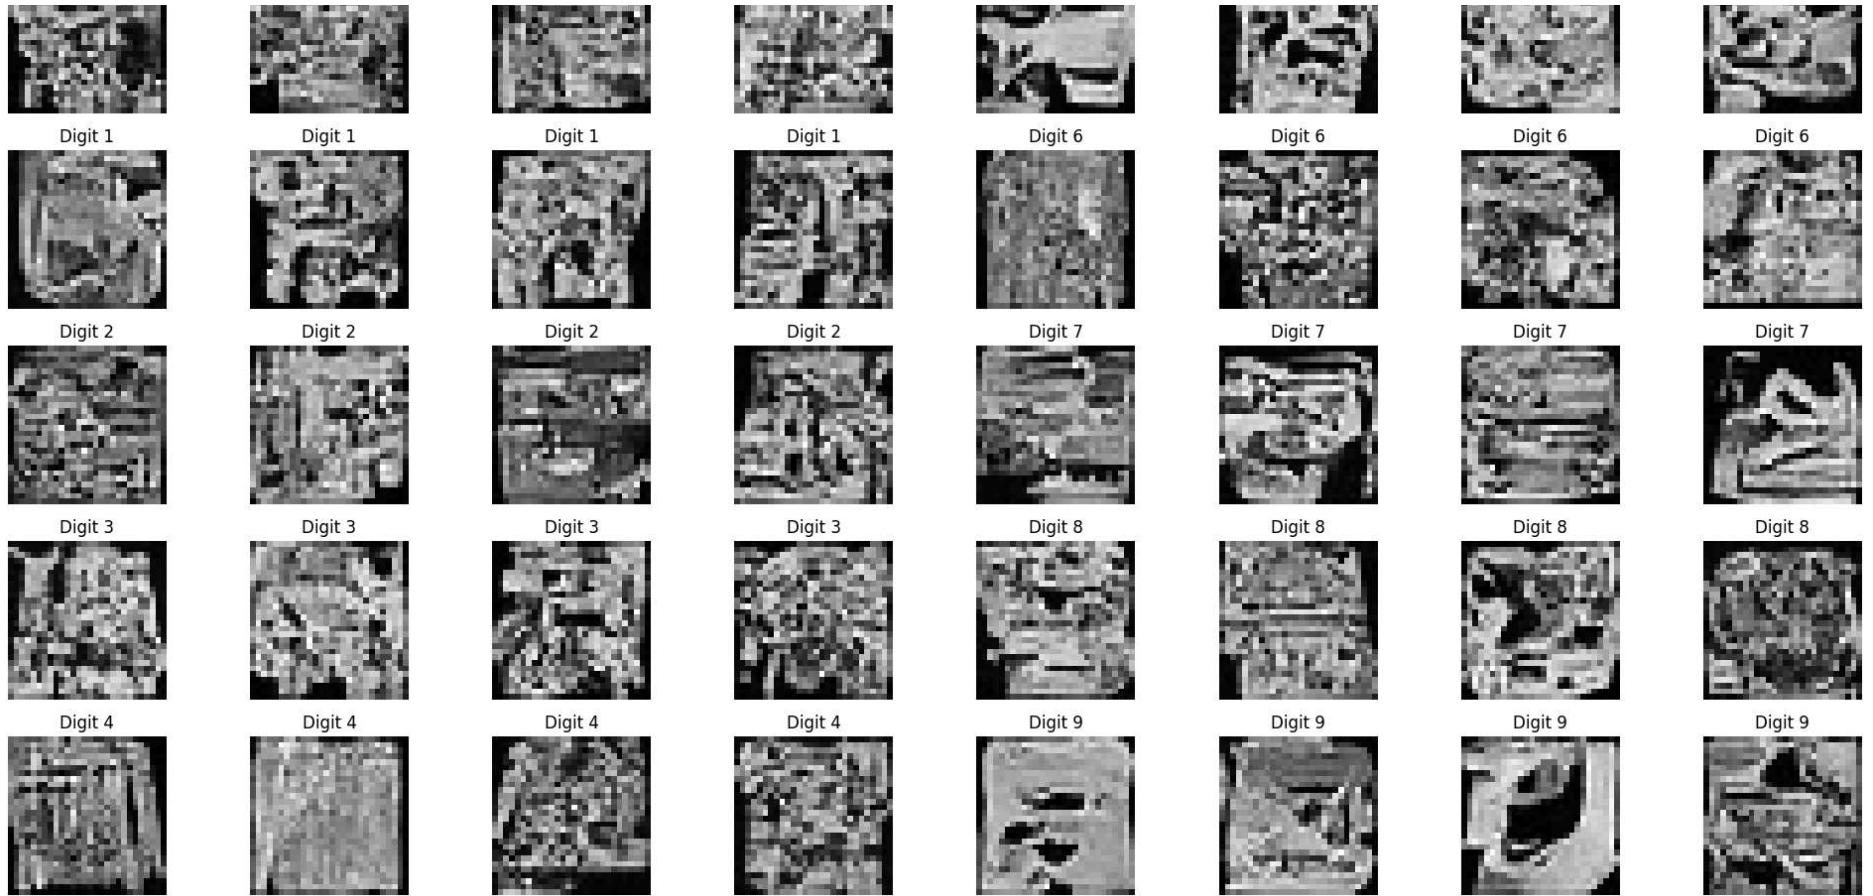
6     model (nn.Module): The trained diffusion model
7     number (int): The digit to generate (0-9)
8     n_samples (int): Number of variations to generate
9
10    Returns:
11        torch.Tensor: Generated images of shape [n_samples, IMG_CH, IMG_SIZE, IMG_SIZE]
12    """
13    model.eval() # Set model to evaluation mode
14    with torch.no_grad(): # No need for gradients during generation
15        # Start with random noise
16        samples = torch.randn(n_samples, IMG_CH, IMG_SIZE, IMG_SIZE).to(device)
17
18        # Set up the number we want to generate
19        c = torch.full((n_samples,), number).to(device)
20        c_one_hot = F.one_hot(c, N_CLASSES).float().to(device)
21        # Correctly sized conditioning mask
22        c_mask = torch.ones_like(c.unsqueeze(-1)).to(device)
23
24        # Display progress information
25        print(f"Generating {n_samples} versions of number {number}...")
26
27        # Remove noise step by step
28        for t in range(n_steps-1, -1, -1):
29            t_batch = torch.full((n_samples,), t).to(device)
30            samples = remove_noise(samples, t_batch, model, c_one_hot, c_mask)
31
32        # Optional: Display occasional progress updates
33        if t % (n_steps // 5) == 0:
34            print(f" Denoising step {n_steps-1-t}/{n_steps-1} completed")
35
36    return samples
37
38 # Generate 4 versions of each number
39 plt.figure(figsize=(20, 10))
40 for i in range(10):
41    # Generate samples for current digit
42    samples = generate_number(model, i, n_samples=4)
43
44    # Display each sample
45    for j in range(4):
46        # Use 2 rows, 10 digits per row, 4 samples per digit
47        # i//5 determines the row (0 or 1)
48        # i%5 determines the position in the row (0-4)
49        # j is the sample index within each digit (0-3)
50        plt.subplot(5, 8, (i%5)*8 + (i//5)*4 + j + 1)
51
52        # Display the image correctly based on channel configuration
53        if IMG_CH == 1: # Grayscale
54            plt.imshow(samples[j][0].cpu(), cmap='gray')
55        else: # Color image
56            img = samples[j].permute(1, 2, 0).cpu()
57            # Rescale from [-1, 1] to [0, 1] if needed
58            if img.min() < 0:
59                img = (img + 1) / 2
60            plt.imshow(img)
61
62        plt.title(f'Digit {i}')
63        plt.axis('off')

```

```
64
65 plt.tight_layout()
66 plt.show()
67
68 # STUDENT ACTIVITY: Try generating the same digit with different noise seeds
69 # This shows the variety of styles the model can produce
70 print("\nSTUDENT ACTIVITY: Generating numbers with different noise seeds")
71
72 # Helper function to generate with seed
73 def generate_with_seed(number, seed_value=42, n_samples=10):
74     torch.manual_seed(seed_value)
75     return generate_number(model, number, n_samples)
76
77 # Pick a image and show many variations
78 # Hint select a image e.g. dog # Change this to any other in the dataset of subset you chose
79 # Hint 2 use variations = generate_with_seed
80 # Hint 3 use plt.figure and plt.imshow to display the variations
81
82 # Enter your code here:
83 digit = 6
84 variations = generate_with_seed(digit, seed_value=42, n_samples=10)
85 plt.figure(figsize=(15, 3))
86 for i in range(10):
87     plt.subplot(1, 10, i+1)
88     plt.imshow(variations[i][0].cpu(), cmap='gray')
89     plt.title(f"Var {i+1}")
90     plt.axis('off')
91 plt.tight_layout()
92 plt.show()
93
```

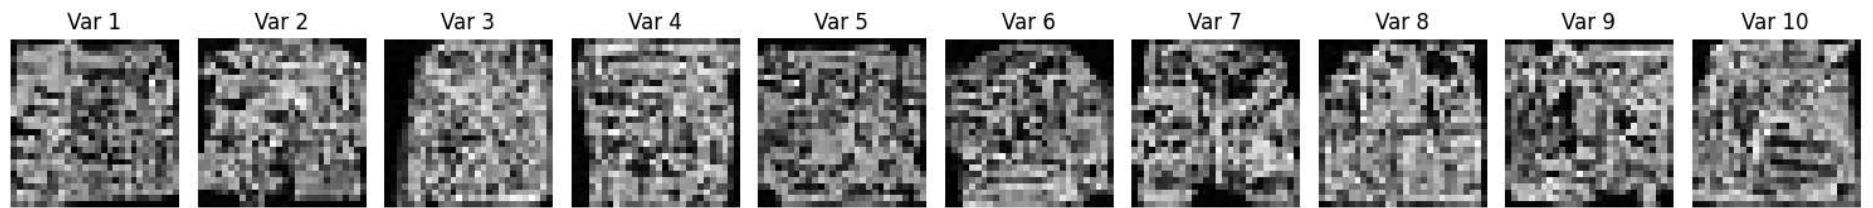
```
Generating 4 versions of number 0...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 1...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 2...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 3...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 4...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 5...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 6...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 7...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 8...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 9...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
```





STUDENT ACTIVITY: Generating numbers with different noise seeds
Generating 10 versions of number 6...

```
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
```



⌄ Step 7: Watching the Generation Process

Let's see how our model turns random noise into clear images, step by step. This helps us understand how the diffusion process works!

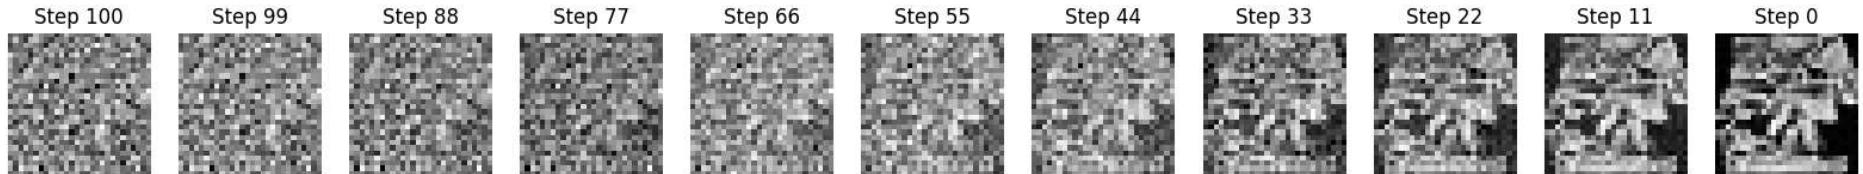
```

1 def visualize_generation_steps(model, number, n_preview_steps=10):
2     """
3         Show how an image evolves from noise to a clear number
4     """
5     model.eval()
6     with torch.no_grad():
7         # Start with random noise
8         x = torch.randn(1, IMG_CH, IMG_SIZE, IMG_SIZE).to(device)
9
10    # Set up which number to generate
11    c = torch.tensor([number]).to(device)
12    c_one_hot = F.one_hot(c, N_CLASSES).float().to(device)
13    c_mask = torch.ones_like(c_one_hot).to(device)
14
15    # Calculate which steps to show
16    steps_to_show = torch.linspace(n_steps-1, 0, n_preview_steps).long()
17
18    # Store images for visualization
19    images = []
20    images.append(x[0].cpu())
21
22    # Remove noise step by step
23    for t in range(n_steps-1, -1, -1):
24        t_batch = torch.full((1,), t).to(device)
25        x = remove_noise(x, t_batch, model, c_one_hot, c_mask)
26
27        if t in steps_to_show:
28            images.append(x[0].cpu())
29
30    # Show the progression
31    plt.figure(figsize=(20, 3))
32    for i, img in enumerate(images):
33        plt.subplot(1, len(images), i+1)
34        if IMG_CH == 1:
35            plt.imshow(img[0], cmap='gray')
36        else:
37            img = img.permute(1, 2, 0)
38            if img.min() < 0:
39                img = (img + 1) / 2
40            plt.imshow(img)
41        step = n_steps if i == 0 else steps_to_show[i-1]
42        plt.title(f'Step {step}')
43        plt.axis('off')
44    plt.show()
45
46 # Show generation process for a few numbers
47 for number in [0, 3, 7]:
48     print(f"\nGenerating number {number}:")
49     visualize_generation_steps(model, number)

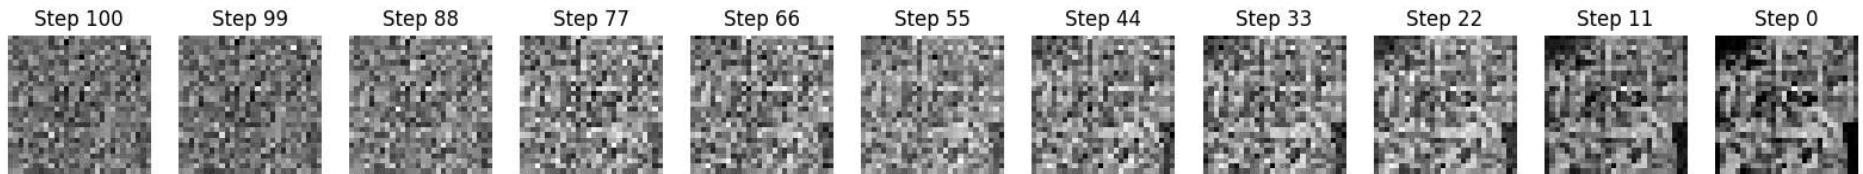
```



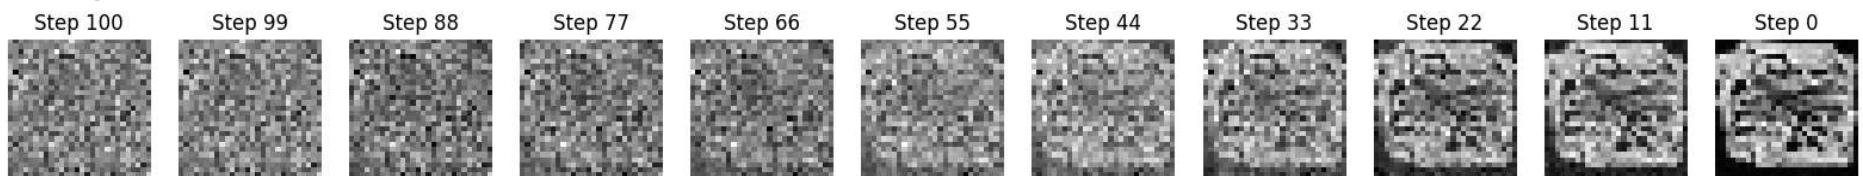
Generating number 0:



Generating number 3:



Generating number 7:



▼ Step 8: Adding CLIP Evaluation

[CLIP](#) is a powerful AI model that can understand both images and text. We'll use it to:

1. Evaluate how realistic our generated images are
2. Score how well they match their intended numbers
3. Help guide the generation process towards better quality

```

1 ## Step 8: Adding CLIP Evaluation
2
3 # CLIP (Contrastive Language-Image Pre-training) is a powerful model by OpenAI that connects text and images.
4 # We'll use it to evaluate how recognizable our generated digits are by measuring how strongly
5 # the CLIP model associates our generated images with text descriptions like "an image of the digit 7".
6
7 # First, we need to install CLIP and its dependencies
8 print("Setting up CLIP (Contrastive Language-Image Pre-training) model...")
9
10 # Track installation status
11 clip_available = False
12
13 try:
14     # Install dependencies first - these help CLIP process text and images
15     print("Installing CLIP dependencies...")

```

```

16 !pip install -q ftfy regex tqdm
17
18 # Install CLIP from GitHub
19 print("Installing CLIP from GitHub repository...")
20 !pip install -q git+https://github.com/openai/CLIP.git
21
22 # Import and verify CLIP is working
23 print("Importing CLIP...")
24 import clip
25
26 # Test that CLIP is functioning
27 models = clip.available_models()
28 print(f"✓ CLIP installation successful! Available models: {models}")
29 clip_available = True
30
31 except ImportError:
32     print("✖ Error importing CLIP. Installation might have failed.")
33     print("Try manually running: !pip install git+https://github.com/openai/CLIP.git")
34     print("If you're in a Colab notebook, try restarting the runtime after installation.")
35
36 except Exception as e:
37     print(f"✖ Error during CLIP setup: {e}")
38     print("Some CLIP functionality may not work correctly.")
39
40 # Provide guidance based on installation result
41 if clip_available:
42     print("\nCLIP is now available for evaluating your generated images!")
43 else:
44     print("\nWARNING: CLIP installation failed. We'll skip the CLIP evaluation parts.")
45
46 # Import necessary libraries
47 import functools
48 import torch.nn.functional as F
49

```

→ Setting up CLIP (Contrastive Language-Image Pre-training) model...
 Installing CLIP dependencies...  44.8/44.8 kB 3.4 MB/s eta 0:00:00
 Installing CLIP from GitHub repository...
 Preparing metadata (setup.py) ... done 
 363.4/363.4 MB 1.1 MB/s eta 0:00:00
 13.8/13.8 MB 64.8 MB/s eta 0:00:00
 24.6/24.6 MB 35.3 MB/s eta 0:00:00
 883.7/883.7 kB 48.4 MB/s eta 0:00:00
 664.8/664.8 MB 1.3 MB/s eta 0:00:00
 211.5/211.5 MB 5.3 MB/s eta 0:00:00
 56.3/56.3 MB 12.9 MB/s eta 0:00:00
 127.9/127.9 MB 8.5 MB/s eta 0:00:00
 207.5/207.5 MB 6.2 MB/s eta 0:00:00
 21.1/21.1 MB 88.3 MB/s eta 0:00:00
 Building wheel for clip (setup.py) ... done
 Importing CLIP...
 ✓ CLIP installation successful! Available models: ['RN50', 'RN101', 'RN50x4', 'RN50x16', 'RN50x64', 'ViT-B/32', 'ViT-B/16', 'ViT-L/14', 'ViT-L/14@336px']
 CLIP is now available for evaluating your generated images!

Below we are creating a helper function to manage GPU memory when using CLIP. CLIP can be memory-intensive, so this will help prevent out-of-memory errors:

```

1 # Memory management decorator to prevent GPU OOM errors
2 def manage_gpu_memory(func):
3     """
4     Decorator that ensures proper GPU memory management.
5
6     This wraps functions that might use large amounts of GPU memory,
7     making sure memory is properly freed after function execution.
8     """
9     @functools.wraps(func)
10    def wrapper(*args, **kwargs):
11        if torch.cuda.is_available():
12            # Clear cache before running function
13            torch.cuda.empty_cache()
14            try:
15                return func(*args, **kwargs)
16            finally:
17                # Clear cache after running function regardless of success/failure
18                torch.cuda.empty_cache()
19        return func(*args, **kwargs)
20    return wrapper

1 =====
2 # Step 8: CLIP Model Loading and Evaluation Setup
3 =====
4 # CLIP (Contrastive Language-Image Pre-training) is a neural network that connects
5 # vision and language. It was trained on 400 million image-text pairs to understand
6 # the relationship between images and their descriptions.
7 # We use it here as an "evaluation judge" to assess our generated images.
8
9 # Load CLIP model with error handling
10 try:
11     # Load the ViT-B/32 CLIP model (Vision Transformer-based)
12     clip_model, clip_preprocess = clip.load("ViT-B/32", device=device)
13     print(f"✓ Successfully loaded CLIP model: {clip_model.visual.__class__.__name__}")
14 except Exception as e:
15     print(f"✗ Failed to load CLIP model: {e}")
16     clip_available = False
17     # Instead of raising an error, we'll continue with degraded functionality
18     print("CLIP evaluation will be skipped. Generated images will still be displayed but without quality scores.")
19
20 def evaluate_with_clip(images, target_number, max_batch_size=16):
21     """
22     Use CLIP to evaluate generated images by measuring how well they match textual descriptions.
23
24     This function acts like an "automatic critic" for our generated digits by measuring:
25     1. How well they match the description of a handwritten digit
26     2. How clear and well-formed they appear to be
27     3. Whether they appear blurry or poorly formed
28
29     The evaluation process works by:
30     - Converting our images to a format CLIP understands
31     - Creating text prompts that describe the qualities we want to measure
32     - Computing similarity scores between images and these text descriptions
33     - Returning normalized scores (probabilities) for each quality
34
35     Args:
36         images (torch.Tensor): Batch of generated images [batch_size, channels, height, width]

```

```

37     target_number (int): The specific digit (0-9) the images should represent
38     max_batch_size (int): Maximum images to process at once (prevents GPU out-of-memory errors)
39
40     Returns:
41         torch.Tensor: Similarity scores tensor of shape [batch_size, 3] with scores for:
42             [good handwritten digit, clear digit, blurry digit]
43             Each row sums to 1.0 (as probabilities)
44     """
45     # If CLIP isn't available, return placeholder scores
46     if not clip_available:
47         print("⚠️ CLIP not available. Returning default scores.")
48         # Equal probabilities (0.33 for each category)
49         return torch.ones(len(images), 3).to(device) / 3
50
51     try:
52         # For large batches, we process in chunks to avoid memory issues
53         # This is crucial when working with big images or many samples
54         if len(images) > max_batch_size:
55             all_similarities = []
56
57             # Process images in manageable chunks
58             for i in range(0, len(images), max_batch_size):
59                 print(f"Processing CLIP batch {i//max_batch_size + 1}/{(len(images)-1)//max_batch_size + 1}")
60                 batch = images[i:i+max_batch_size]
61
62                 # Use context managers for efficiency and memory management:
63                 # - torch.no_grad(): disables gradient tracking (not needed for evaluation)
64                 # - torch.cuda.amp.autocast(): uses mixed precision to reduce memory usage
65                 with torch.no_grad(), torch.cuda.amp.autocast():
66                     batch_similarities = _process_clip_batch(batch, target_number)
67                     all_similarities.append(batch_similarities)
68
69                 # Explicitly free GPU memory between batches
70                 # This helps prevent cumulative memory buildup that could cause crashes
71                 torch.cuda.empty_cache()
72
73             # Combine results from all batches into a single tensor
74             return torch.cat(all_similarities, dim=0)
75         else:
76             # For small batches, process all at once
77             with torch.no_grad(), torch.cuda.amp.autocast():
78                 return _process_clip_batch(images, target_number)
79
80     except Exception as e:
81         # If anything goes wrong, log the error but don't crash
82         print(f"❌ Error in CLIP evaluation: {e}")
83         print(f"Traceback: {traceback.format_exc()}")
84         # Return default scores so the rest of the notebook can continue
85         return torch.ones(len(images), 3).to(device) / 3
86
87 def _process_clip_batch(images, target_number):
88     """
89     Core CLIP processing function that computes similarity between images and text descriptions.
90
91     This function handles the technical details of:
92     1. Preparing relevant text prompts for evaluation
93     2. Preprocessing images to CLIP's required format
94     3. Extracting feature embeddings from both images and text

```

```

95    4. Computing similarity scores between these embeddings
96
97    The function includes advanced error handling for GPU memory issues,
98    automatically reducing batch size if out-of-memory errors occur.
99
100   Args:
101       images (torch.Tensor): Batch of images to evaluate
102       target_number (int): The digit these images should represent
103
104   Returns:
105       torch.Tensor: Normalized similarity scores between images and text descriptions
106   """
107   try:
108       # Create text descriptions (prompts) to evaluate our generated digits
109       # We check three distinct qualities:
110       # 1. If it looks like a handwritten example of the target digit
111       # 2. If it appears clear and well-formed
112       # 3. If it appears blurry or poorly formed (negative case)
113       text_inputs = torch.cat([
114           clip.tokenize(f"A handwritten number {target_number}"),
115           clip.tokenize(f"A clear, well-written digit {target_number}"),
116           clip.tokenize(f"A blurry or unclear number")
117       ]).to(device)
118
119       # Process images for CLIP, which requires specific formatting:
120
121       # 1. Handle different channel configurations (dataset-dependent)
122       if IMG_CH == 1:
123           # CLIP expects RGB images, so we repeat the grayscale channel 3 times
124           # For example, MNIST/Fashion-MNIST are grayscale (1-channel)
125           images_rgb = images.repeat(1, 3, 1, 1)
126       else:
127           # For RGB datasets like CIFAR-10/CelebA, we can use as-is
128           images_rgb = images
129
130       # 2. Normalize pixel values to [0,1] range if needed
131       # Different datasets may have different normalization ranges
132       if images_rgb.min() < 0: # If normalized to [-1,1] range
133           images_rgb = (images_rgb + 1) / 2 # Convert to [0,1] range
134
135       # 3. Resize images to CLIP's expected input size (224x224 pixels)
136       # CLIP was trained on this specific resolution
137       resized_images = F.interpolate(images_rgb, size=(224, 224),
138                                      mode='bilinear', align_corners=False)
139
140       # Extract feature embeddings from both images and text prompts
141       # These are high-dimensional vectors representing the content
142       image_features = clip_model.encode_image(resized_images)
143       text_features = clip_model.encode_text(text_inputs)
144
145       # Normalize feature vectors to unit length (for cosine similarity)
146       # This ensures we're measuring direction, not magnitude
147       image_features = image_features / image_features.norm(dim=-1, keepdim=True)
148       text_features = text_features / text_features.norm(dim=-1, keepdim=True)
149
150       # Calculate similarity scores between image and text features
151       # The matrix multiplication computes all pairwise dot products at once
152       # Multiplying by 100 scales to percentage-like values before applying softmax

```

```
153     similarity = (100.0 * image_features @ text_features.T).softmax(dim=-1)
154
155     return similarity
156
157 except RuntimeError as e:
158     # Special handling for CUDA out-of-memory errors
159     if "out of memory" in str(e):
160         # Free GPU memory immediately
161         torch.cuda.empty_cache()
162
163     # If we're already at batch size 1, we can't reduce further
164     if len(images) <= 1:
165         print("✖ Out of memory even with batch size 1. Cannot process.")
166         return torch.ones(len(images), 3).to(device) / 3
167
168     # Adaptive batch size reduction - recursively try with smaller batches
169     # This is an advanced technique to handle limited GPU memory gracefully
170     half_size = len(images) // 2
171     print(f"⚠️ Out of memory. Reducing batch size to {half_size}.")
172
173     # Process each half separately and combine results
174     # This recursive approach will keep splitting until processing succeeds
175     first_half = _process_clip_batch(images[:half_size], target_number)
176     second_half = _process_clip_batch(images[half_size:], target_number)
177
178     # Combine results from both halves
179     return torch.cat([first_half, second_half], dim=0)
180
181     # For other errors, propagate upward
182     raise e
183
184 #####CLIP Evaluation - Generate and Analyze Sample Digits#####
185 # CLIP Evaluation - Generate and Analyze Sample Digits
186 #####
187 # This section demonstrates how to use CLIP to evaluate generated digits
188 # We'll generate examples of all ten digits and visualize the quality scores
189
190 try:
191     for number in range(10):
192         print(f"\nGenerating and evaluating number {number}...")
193
194     # Generate 4 different variations of the current digit
195     samples = generate_number(model, number, n_samples=4)
196
197     # Evaluate quality with CLIP (without tracking gradients for efficiency)
198     with torch.no_grad():
199         similarities = evaluate_with_clip(samples, number)
200
201     # Create a figure to display results
202     plt.figure(figsize=(15, 3))
203
204     # Show each sample with its CLIP quality scores
205     for i in range(4):
206         plt.subplot(1, 4, i+1)
207
208         # Display the image with appropriate formatting based on dataset type
209         if IMG_CH == 1: # Grayscale images (MNIST, Fashion-MNIST)
210             plt.imshow(samples[i][0].cpu(), cmap='gray')
```

```

211         else: # Color images (CIFAR-10, CelebA)
212             img = samples[i].permute(1, 2, 0).cpu() # Change format for matplotlib
213             if img.min() < 0: # Handle [-1,1] normalization
214                 img = (img + 1) / 2 # Convert to [0,1] range
215             plt.imshow(img)
216
217     # Extract individual quality scores for display
218     # These represent how confidently CLIP associates the image with each description
219     good_score = similarities[i][0].item() * 100 # Handwritten quality
220     clear_score = similarities[i][1].item() * 100 # Clarity quality
221     blur_score = similarities[i][2].item() * 100 # Blurriness assessment
222
223     # Color-code the title based on highest score category:
224     # - Green: if either "good handwritten" or "clear" score is highest
225     # - Red: if "blurry" score is highest (poor quality)
226     max_score_idx = torch.argmax(similarities[i]).item()
227     title_color = 'green' if max_score_idx < 2 else 'red'
228
229     # Show scores in the plot title
230     plt.title(f'Number {number}\nGood: {good_score:.0f}%\nClear: {clear_score:.0f}%\nBlurry: {blur_score:.0f}%',
231             color=title_color)
232     plt.axis('off')
233
234     plt.tight_layout()
235     plt.show()
236     plt.close() # Properly close figure to prevent memory leaks
237
238     # Clean up GPU memory after processing each number
239     # This is especially important for resource-constrained environments
240     torch.cuda.empty_cache()
241
242 except Exception as e:
243     # Comprehensive error handling to help students debug issues
244     print(f"✖ Error in generation and evaluation loop: {e}")
245     print("Detailed error information:")
246     import traceback
247     traceback.print_exc()
248
249     # Clean up resources even if we encounter an error
250     if torch.cuda.is_available():
251         print("Clearing GPU cache...")
252         torch.cuda.empty_cache()
253
254 =====
255 # STUDENT ACTIVITY: Exploring CLIP Evaluation
256 =====
257 # This section provides code templates for students to experiment with
258 # evaluating larger batches of generated digits using CLIP.
259
260 print("\nSTUDENT ACTIVITY:")
261 print("Try the code below to evaluate a larger sample of a specific digit")
262 print("""
263 # Example: Generate and evaluate 10 examples of the digit 6
264 # digit = 6
265 # samples = generate_number(model, digit, n_samples=10)
266 # similarities = evaluate_with_clip(samples, digit)
267 #
268 # # Calculate what percentage of samples CLIP considers "good quality"

```

```
269 # # (either "good handwritten" or "clear" score exceeds "blurry" score)
270 # good_or_clear = (similarities[:,0] + similarities[:,1] > similarities[:,2]).float().mean()
271 # print(f"CLIP recognized {good_or_clear.item() * 100:.1f}% of the digits as good examples of {digit}")
272 #
273 # # Display a grid of samples with their quality scores
274 # plt.figure(figsize=(15, 8))
275 # for i in range(len(samples)):
276 #     plt.subplot(2, 5, i+1)
277 #     plt.imshow(samples[i][0].cpu(), cmap='gray')
278 #     quality = "Good" if similarities[i,0] + similarities[i,1] > similarities[i,2] else "Poor"
279 #     plt.title(f"Sample {i+1}: {quality}", color='green' if quality == "Good" else 'red')
280 #     plt.axis('off')
281 # plt.tight_layout()
282 # plt.show()
283 """
284 # Enter your code here:
285 # Example: Generate and evaluate 10 examples of the digit 6 (trouser in Fashion-MNIST)
286 digit = 6 # Class 6 in Fashion-MNIST is "shirt"
287 samples = generate_number(model, digit, n_samples=10)
288 similarities = evaluate_with_clip(samples, digit)
289
290 # Calculate what percentage of samples CLIP considers "good quality"
291 # (either "good handwritten" or "clear" score exceeds "blurry" score)
292 good_or_clear = (similarities[:,0] + similarities[:,1] > similarities[:,2]).float().mean()
293 print(f"CLIP recognized {good_or_clear.item() * 100:.1f}% of the samples as good examples of class {digit}")
294
295 # Display a grid of samples with their quality scores
296 plt.figure(figsize=(15, 8))
297 for i in range(len(samples)):
298     plt.subplot(2, 5, i+1)
299     plt.imshow(samples[i][0].cpu(), cmap='gray')
300     quality = "Good" if similarities[i,0] + similarities[i,1] > similarities[i,2] else "Poor"
301     plt.title(f"Sample {i+1}: {quality}", color='green' if quality == "Good" else 'red')
302     plt.axis('off')
303 plt.tight_layout()
304 plt.show()
```

100% | 338M/338M [00:09<00:00, 38.3MiB/s]
✓ Successfully loaded CLIP model: VisionTransformer

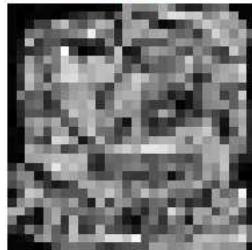
Generating and evaluating number 0...

Generating 4 versions of number 0...

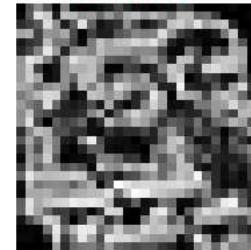
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed

<ipython-input-27-4799bd7c2ed7>:77: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with torch.no_grad(), torch.cuda.amp.autocast():

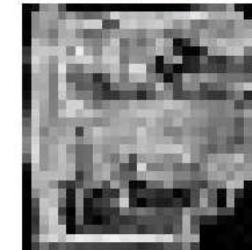
Number 0
Good: 2%
Clear: 73%
Blurry: 25%



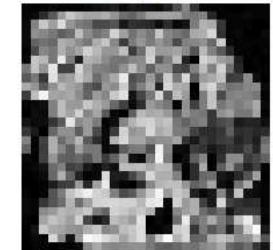
Number 0
Good: 2%
Clear: 33%
Blurry: 65%



Number 0
Good: 2%
Clear: 64%
Blurry: 34%



Number 0
Good: 3%
Clear: 53%
Blurry: 44%

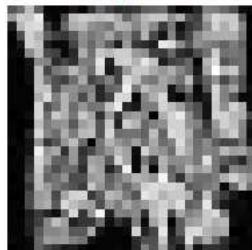


Generating and evaluating number 1...

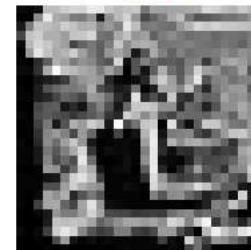
Generating 4 versions of number 1...

Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed

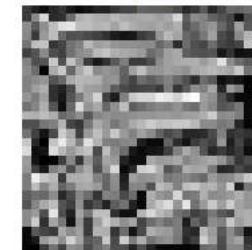
Number 1
Good: 1%
Clear: 63%
Blurry: 36%



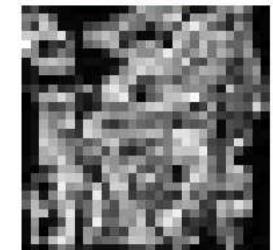
Number 1
Good: 1%
Clear: 62%
Blurry: 37%



Number 1
Good: 1%
Clear: 64%
Blurry: 35%



Number 1
Good: 1%
Clear: 74%
Blurry: 25%



Generating and evaluating number 2...

Generating 4 versions of number 2...

Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed

Number 2
Good: 1%

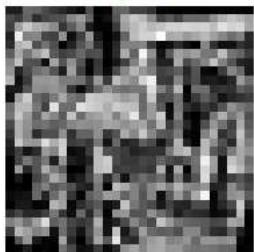
https://colab.research.google.com/drive/1PZjBlBoSnfH7NqGkVA_pqbqo2Dh3TRY#scrollTo=YR_xQNAQg-eO&printMode=true

Number 2
Good: 1%

Number 2
Good: 1%

Number 2
Good: 0%

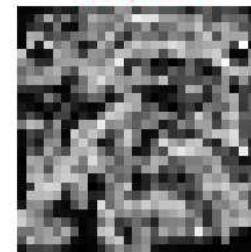
Good: 1%
Clear: 55%
Blurry: 44%



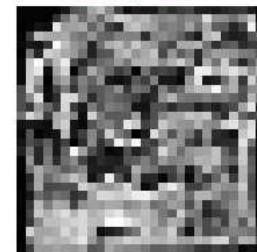
Good: 1%
Clear: 35%
Blurry: 64%



Good: 1%
Clear: 40%
Blurry: 59%



Good: 0%
Clear: 67%
Blurry: 33%

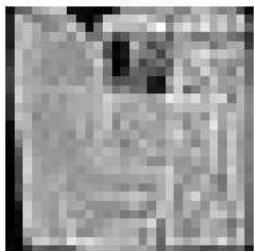


Generating and evaluating number 3...

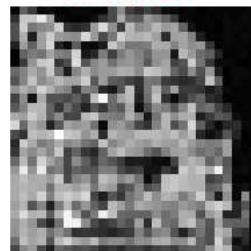
Generating 4 versions of number 3...

Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed

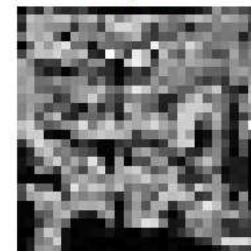
Number 3
Good: 1%
Clear: 22%
Blurry: 77%



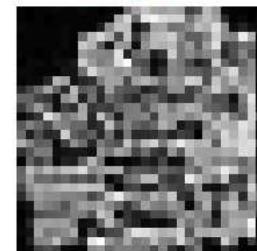
Number 3
Good: 1%
Clear: 43%
Blurry: 56%



Number 3
Good: 2%
Clear: 59%
Blurry: 38%



Number 3
Good: 2%
Clear: 58%
Blurry: 40%

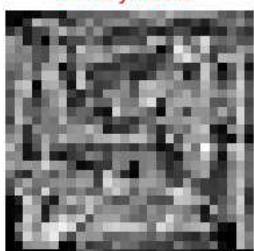


Generating and evaluating number 4...

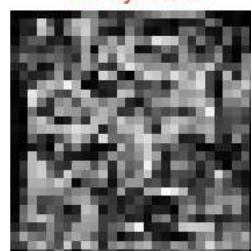
Generating 4 versions of number 4...

Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed

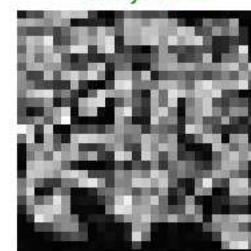
Number 4
Good: 3%
Clear: 35%
Blurry: 63%



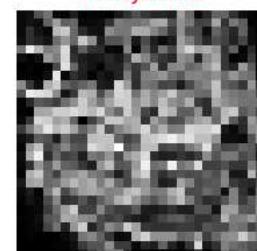
Number 4
Good: 2%
Clear: 20%
Blurry: 78%



Number 4
Good: 4%
Clear: 51%
Blurry: 44%



Number 4
Good: 2%
Clear: 38%
Blurry: 60%



Generating and evaluating number 5...

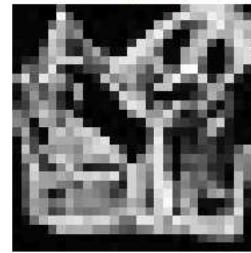
Generating 4 versions of number 5...

Denoising step 19/99 completed
 Denoising step 39/99 completed
 Denoising step 59/99 completed
 Denoising step 79/99 completed
 Denoising step 99/99 completed

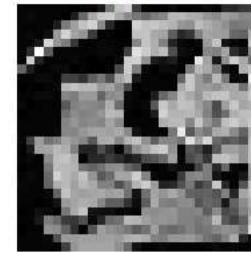
Number 5
 Good: 3%
 Clear: 33%
 Blurry: 64%



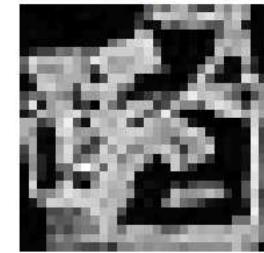
Number 5
 Good: 1%
 Clear: 78%
 Blurry: 21%



Number 5
 Good: 2%
 Clear: 69%
 Blurry: 29%



Number 5
 Good: 1%
 Clear: 50%
 Blurry: 49%

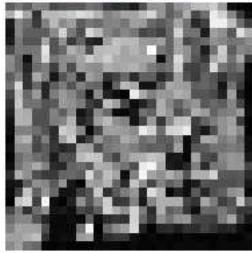


Generating and evaluating number 6...

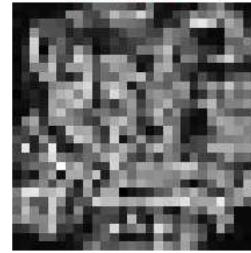
Generating 4 versions of number 6...

Denoising step 19/99 completed
 Denoising step 39/99 completed
 Denoising step 59/99 completed
 Denoising step 79/99 completed
 Denoising step 99/99 completed

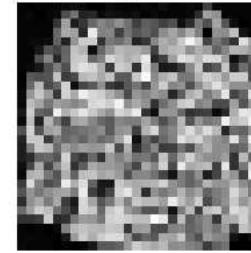
Number 6
 Good: 3%
 Clear: 62%
 Blurry: 36%



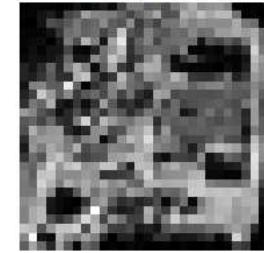
Number 6
 Good: 1%
 Clear: 34%
 Blurry: 65%



Number 6
 Good: 4%
 Clear: 51%
 Blurry: 46%



Number 6
 Good: 1%
 Clear: 59%
 Blurry: 40%



Generating and evaluating number 7...

Generating 4 versions of number 7...

Denoising step 19/99 completed
 Denoising step 39/99 completed
 Denoising step 59/99 completed
 Denoising step 79/99 completed
 Denoising step 99/99 completed

Number 7
 Good: 2%
 Clear: 33%
 Blurry: 65%



Number 7
 Good: 1%
 Clear: 15%
 Blurry: 84%



Number 7
 Good: 2%
 Clear: 57%
 Blurry: 40%



Number 7
 Good: 2%
 Clear: 54%
 Blurry: 45%

